

CCCN 312 Computer Networks

Chapter 4

Network Layer: Data Plane

A note on the use of these PowerPoint slides:


We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

For a revision history, see the slide note for this page.

Thanks and enjoy! JFK/KWR

All material copyright 1996-2020
J.F Kurose and K.W. Ross, All Rights Reserved



*Computer Networking: A
Top-Down Approach*

8th edition

Jim Kurose, Keith Ross
Pearson, 2020

Outline

1. Introduction



2. Application layer



3. Transport layer



4. Network layer: Data Plane

5. Network layer: Control Plane

6. Link layer

Network layer: our goals

- understand principles behind network layer services, focusing on data plane:
 - network layer service models
 - forwarding versus routing
 - how a router works
 - addressing
 - forwarding
 - Internet architecture
- instantiation, implementation in the Internet
 - IP protocol
 - NAT, DHCP


Network layer: “data plane” roadmap

- Network layer: overview
 - data plane
 - control plane
- What's inside a router
 - input ports, switching, output ports
 - buffer management, scheduling
- IP: the Internet Protocol
 - datagram format
 - addressing
 - network address translation
 - IPv6



Network-layer services and protocols

- transport segment from sending to receiving host
 - sender: encapsulates segments into datagrams, passes to link layer
 - receiver: delivers segments to transport layer protocol
- network layer protocols in *every Internet device*: hosts, routers
- routers:
 - examines header fields in all IP datagrams passing through it
 - moves datagrams from input ports to output ports to transfer datagrams along end-end path



Two key network-layer functions

network-layer functions:

- 1 ■ **forwarding**: move packets from a router's input link to appropriate router output link
- 2 ■ **routing**: determine route taken by packets from source to destination
 - *routing algorithms*

analogy: taking a trip

- **forwarding**: process of getting through single interchange
- **routing**: process of planning trip from source to destination

تنفيذ
الخطة

الخطة



forwarding




routing

Network layer: data plane, control plane

Data plane:

- *local*, per-router function
- determines how datagram arriving on router input port is *forwarded* to router output port



Control plane

- *network-wide* logic global
- determines how datagram is *routed* among routers along end-end path from source host to destination host
- two control-plane approaches:
 1. *traditional routing algorithms*: implemented in routers
 2. *software-defined networking (SDN)*: implemented in (remote) servers

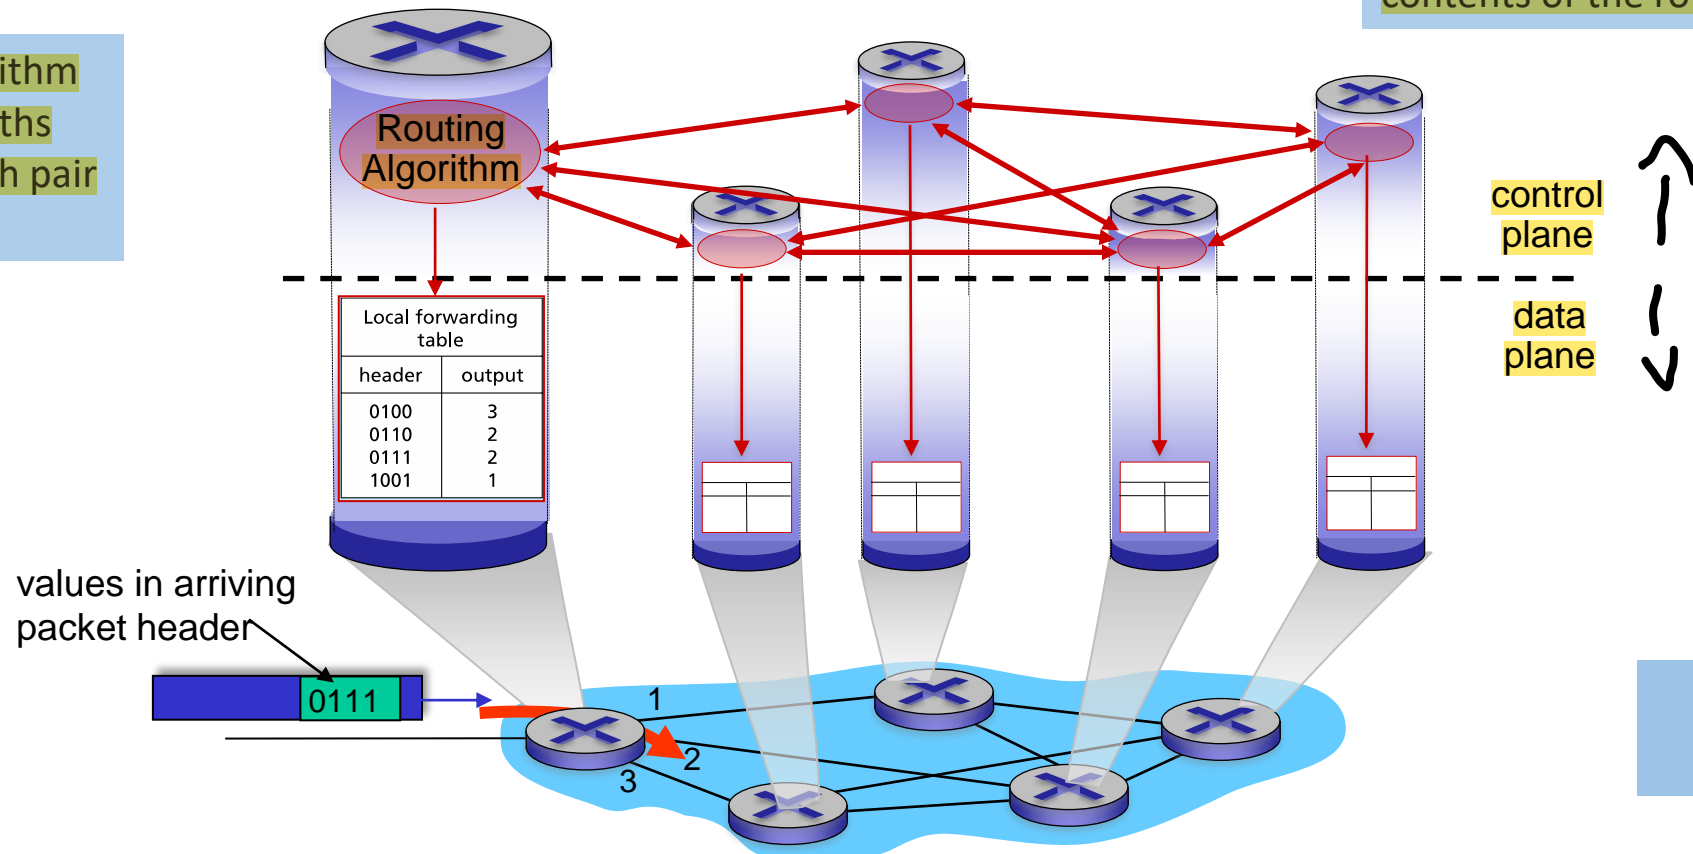
Per-router control plane

Individual routing algorithm components *in each and every router* interact in the control plane

(interact by exchanging routing messages)

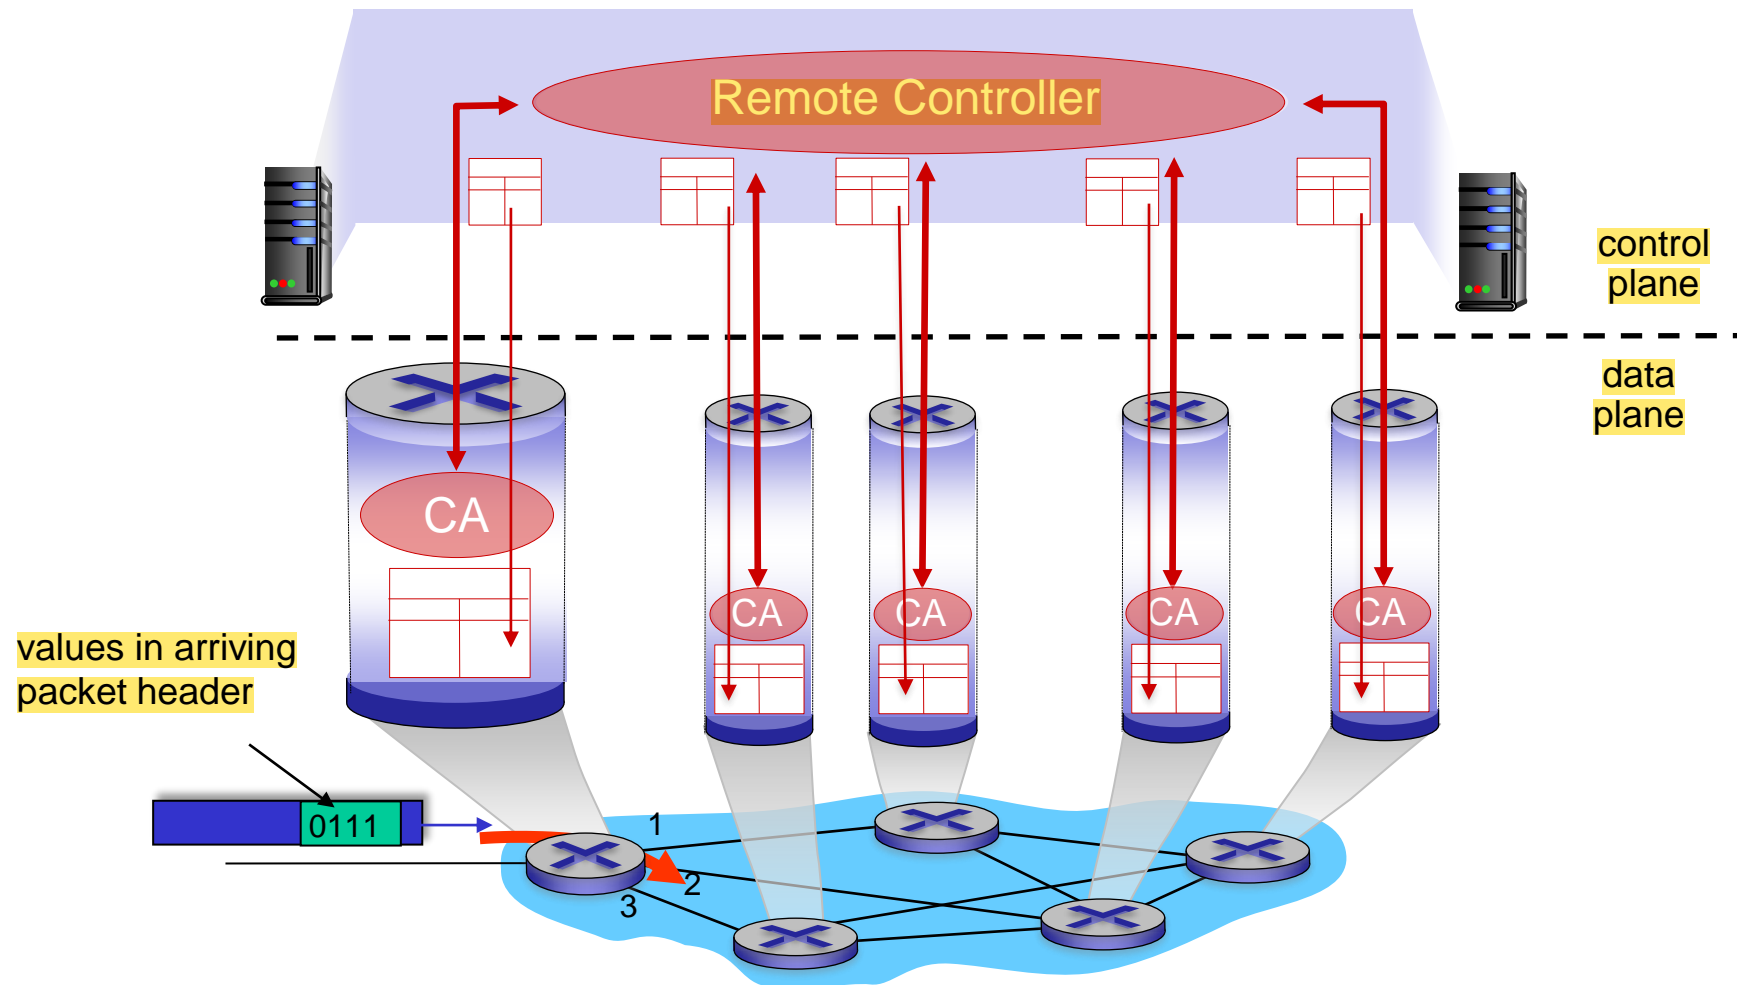
routing algorithm
computes paths
between each pair
of routers

the routing algorithm determines the
contents of the routers' forwarding tables



Software-Defined Networking (SDN) control plane

Remote controller computes, installs forwarding tables in routers



2

This is the concept of SDN

Network service model

Q: What *service model* for “channel” transporting datagrams from sender to receiver?

example services for
individual datagrams:

- guaranteed delivery
- guaranteed delivery with less than 40 msec delay

example services for a *flow* of datagrams:

- in-order datagram delivery
- guaranteed minimum bandwidth to flow
- restrictions on changes in inter-packet spacing

Network-layer service model

Network Architecture	Service Model	Quality of Service (QoS) Guarantees ?			
		Bandwidth	Loss	Order	Timing
Internet	best effort	none	no	no	no

Internet “best effort” service model

No guarantees on:

- i. successful datagram delivery to destination
- ii. timing or order of delivery
- iii. bandwidth available to end-end flow

Network-layer service model

FYI

Network Architecture	Service Model	Quality of Service (QoS) Guarantees ?			
		Bandwidth	Loss	Order	Timing
Internet	best effort	none	no	no	no
ATM	Constant Bit Rate	Constant rate	yes	yes	yes
ATM	Available Bit Rate	Guaranteed min	no	yes	no
Internet	Intserv Guaranteed (RFC 1633)	yes	yes	yes	yes
Internet	Diffserv (RFC 2475)	possible	possibly	possibly	no

**There are other service models with better QoS,
but none of them succeeded**

Reflections on best-effort service:



- **simplicity of mechanism** has allowed Internet to be widely deployed adopted
- sufficient **provisioning of bandwidth** allows performance of real-time applications (e.g., interactive voice, video) to be “good enough” for “most of the time”
- **replicated, application-layer distributed services** (datacenters, content distribution networks) connecting close to clients’ networks, allow services to be provided from multiple locations
- congestion control of “elastic” services helps

It's hard to argue with success of best-effort service model

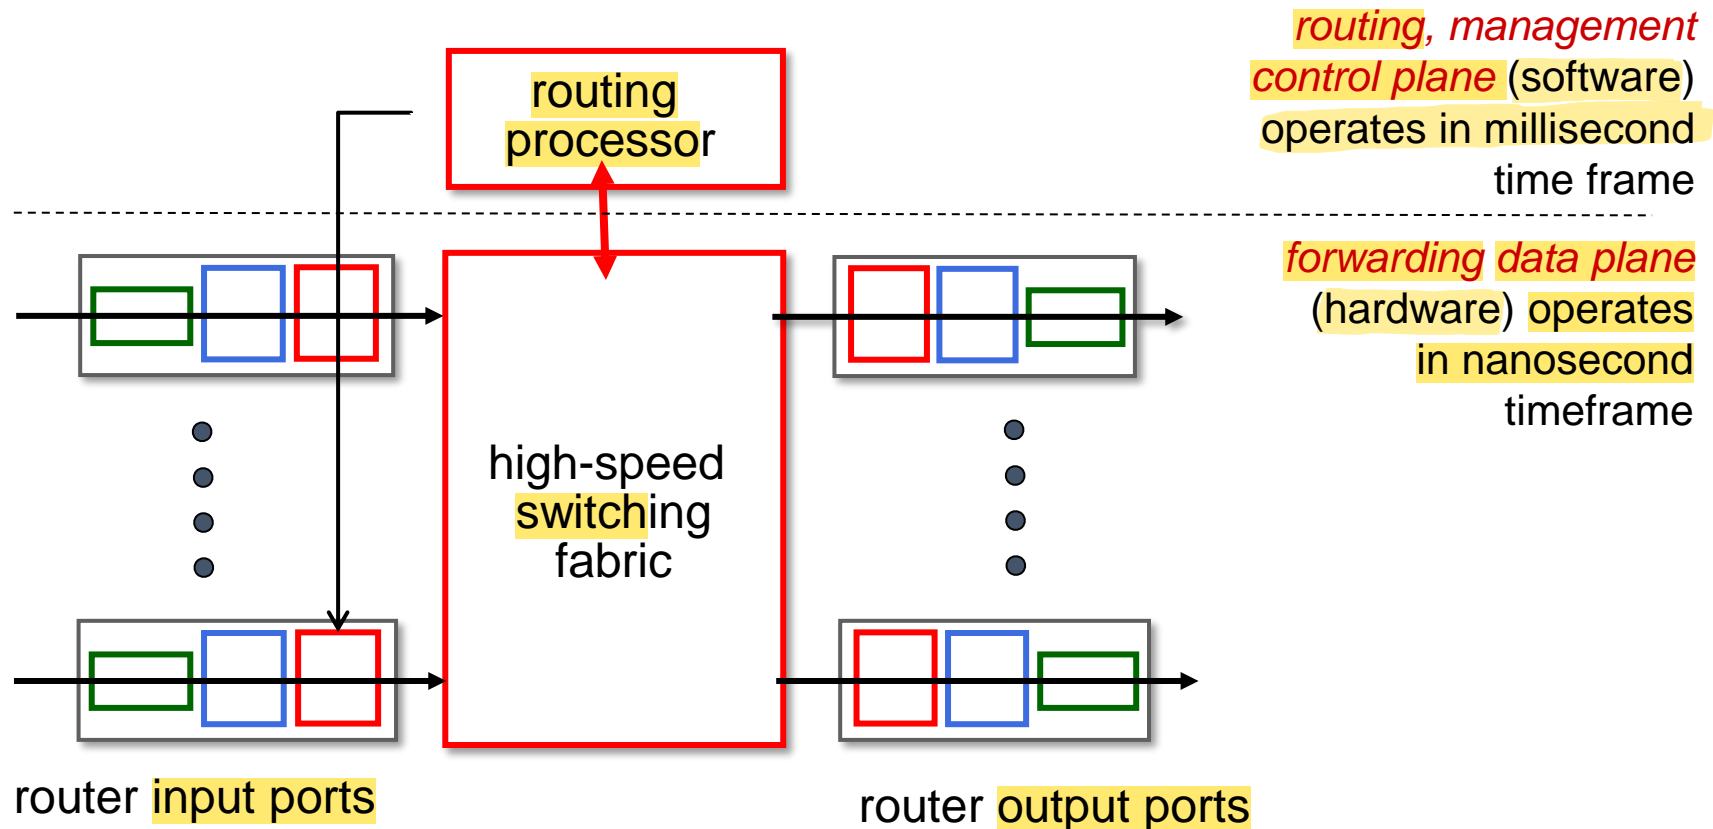
Network layer: “data plane” roadmap

- Network layer: overview
 - data plane
 - control plane
- What’s inside a router
 - input ports, switching, output ports
 - buffer management, scheduling
- IP: the Internet Protocol
 - datagram format
 - addressing
 - network address translation
 - IPv6

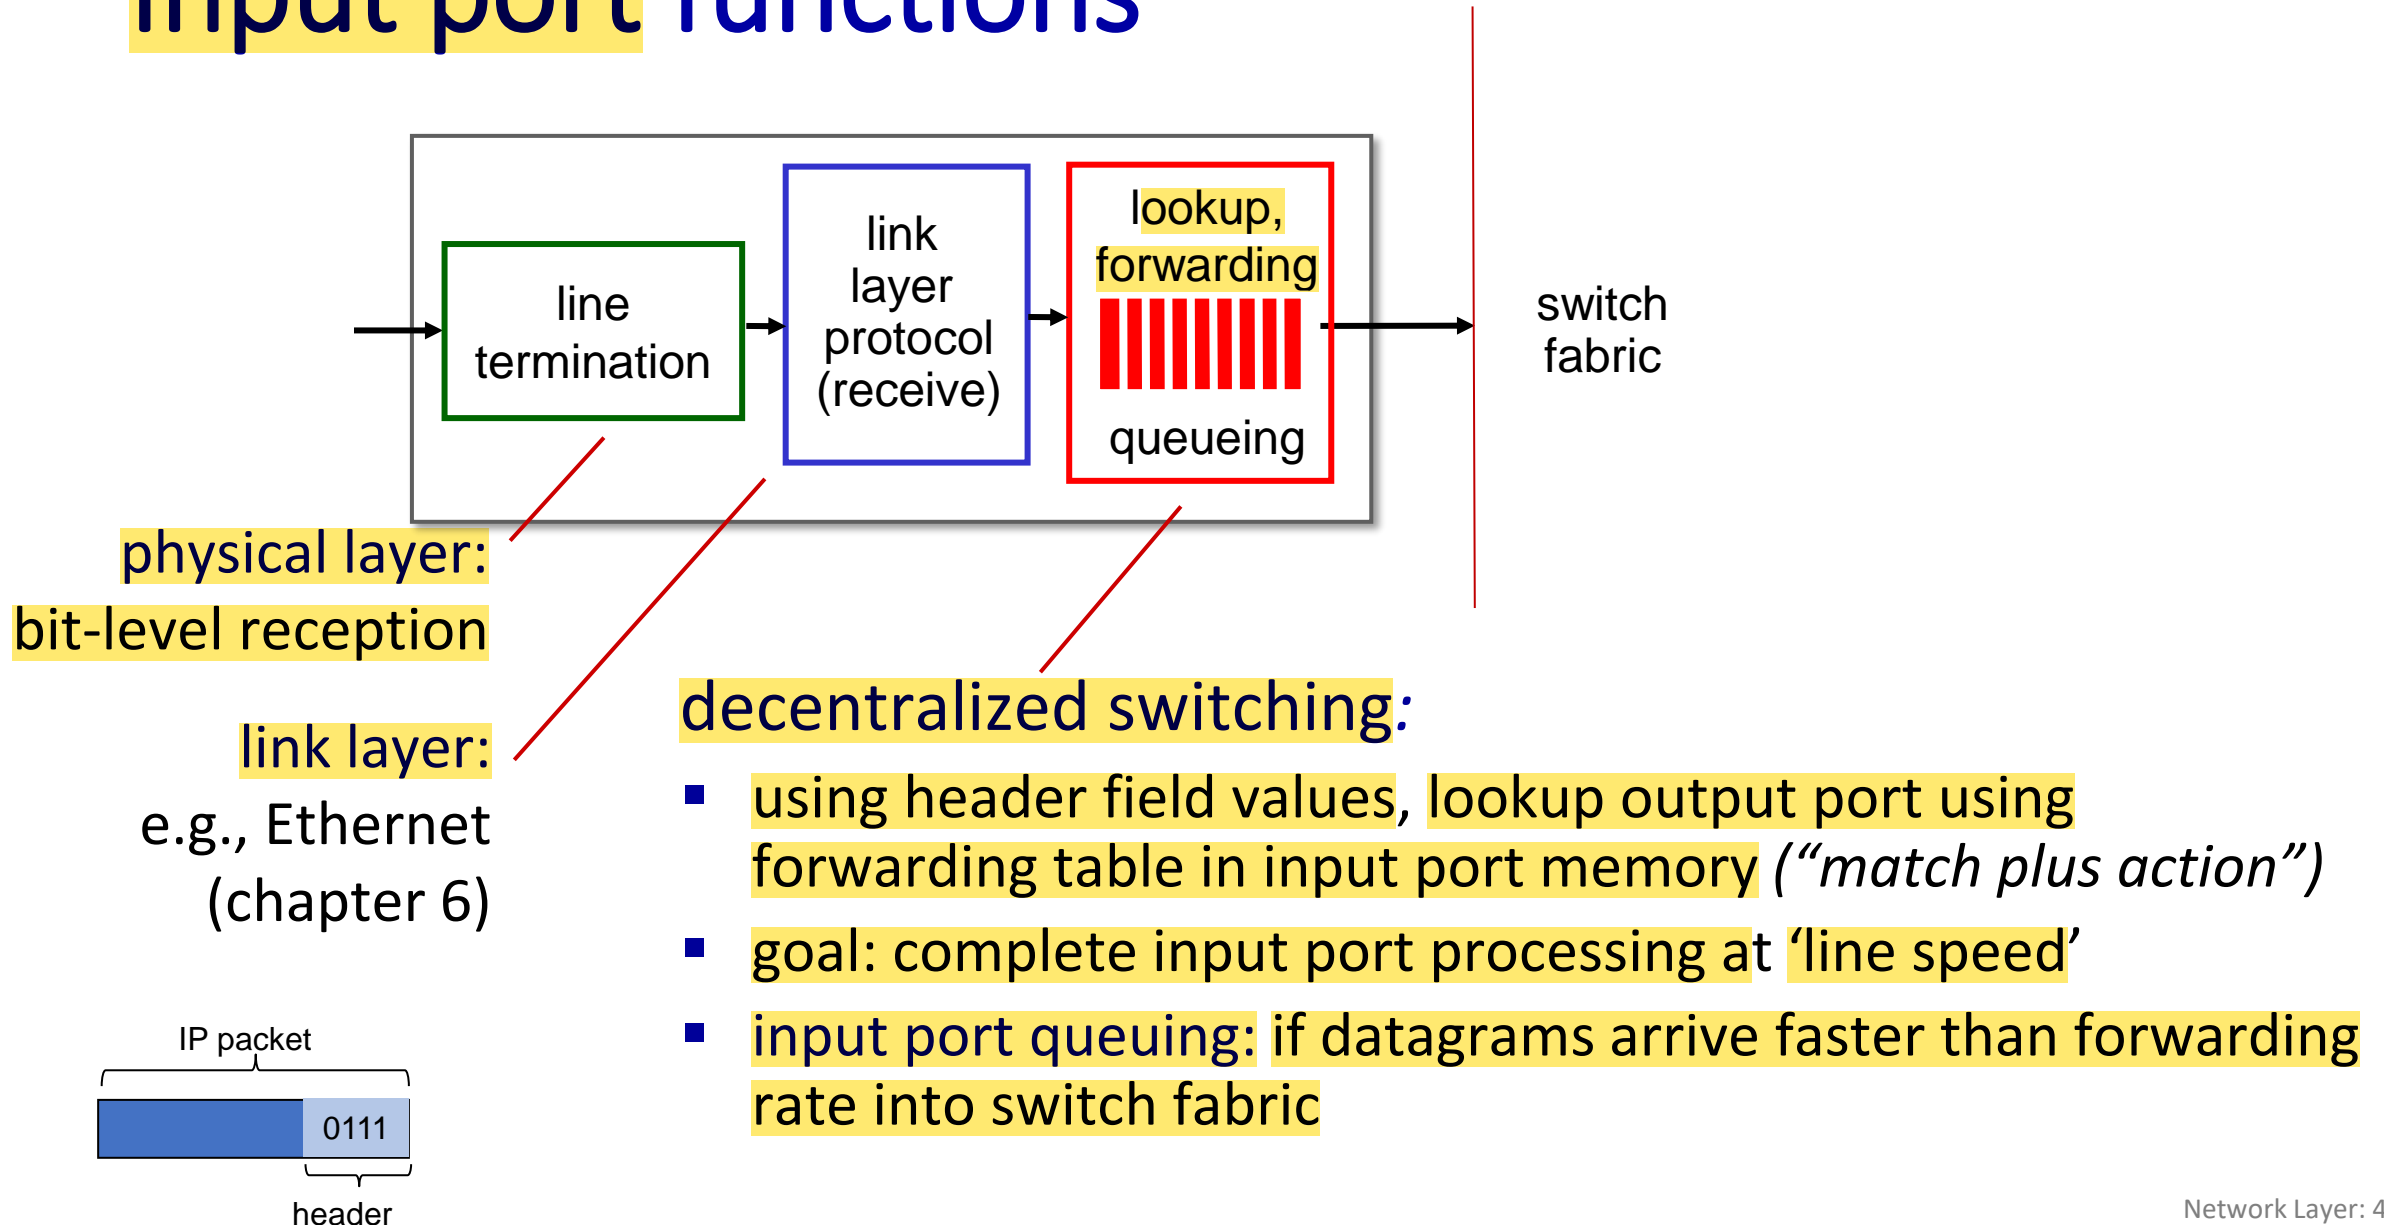


Router architecture overview

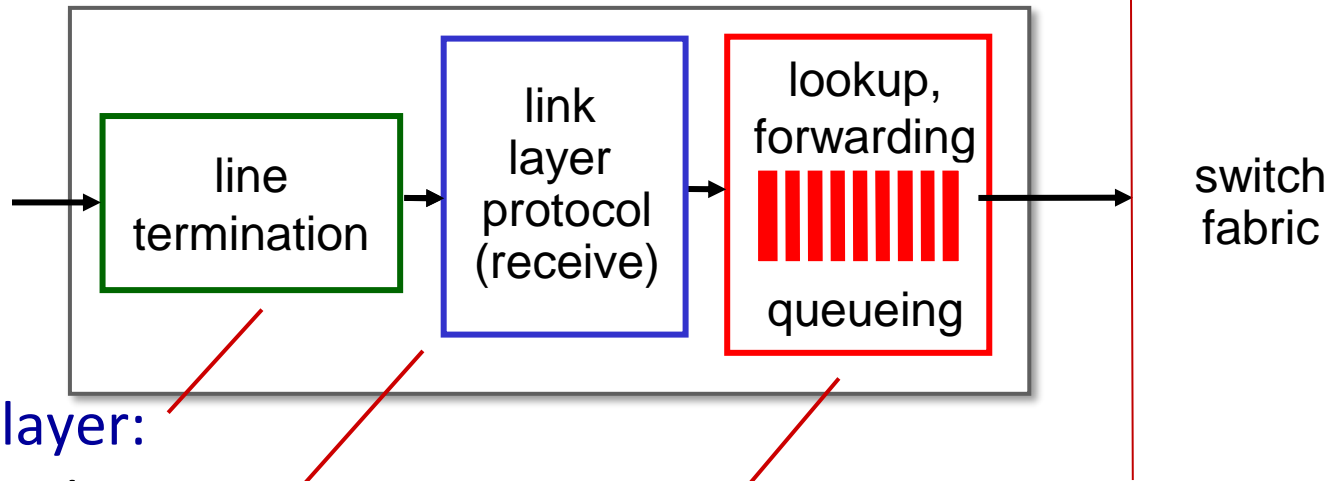
high-level view of generic router architecture:



Input port functions



Input port functions



physical layer:
bit-level reception

link layer:
e.g., Ethernet
(chapter 6)

decentralized switching:

- using header field values, lookup output port using forwarding table in input port memory (*"match plus action"*)
- 1 ■ **destination-based forwarding:** forward based only on destination IP address (traditional)
- 2 ■ **generalized forwarding:** forward based on any set of header field values

Destination-based forwarding

<i>forwarding table</i>	
Destination Address Range	Link Interface
11001000 00010111 00010000 00000000 through 11001000 00010111 00010000 00000100 through 11001000 00010111 00010000 00000111	n 3
11001000 00010111 00011000 11111111 through 11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	2
otherwise	3

Q: but what happens if ranges don't divide up so nicely?

Longest prefix matching

longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

examples:

11001000 00010111 00010110 10100001

which interface? 0

11001000 00010111 00011000 10101010

which interface? 1

Longest prefix matching

longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range					Link interface
11001000	00010111	00010***	*****		0
11001000	00010111	00011000	*****		1
11001000	match! 1	00011***	*****		2
otherwise					3

examples:

11001000	00010111	00010110	10100001	which interface?
11001000	00010111	00011000	10101010	which interface?

Longest prefix matching

longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range				Link interface
11001000	00010111	00010***	*****	0
11001000	00010111	00011000	*****	1
11001000	00010111	00011***	*****	2
otherwise				3

match!

which interface?

which interface?

examples:

11001000	00010111	00010110	10100001
11001000	00010111	00011000	10101010

Longest prefix matching

longest prefix match

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range				Link interface
11001000	00010111	00010***	*****	0
11001000	00010111	00011000	*****	1
11001000	00010111	00011***	*****	2
otherwise				3
11001000	00010111	00010110	10100001	which interface?
11001000	00010111	00011000	10101010	which interface?

match! ✓ longest

examples:

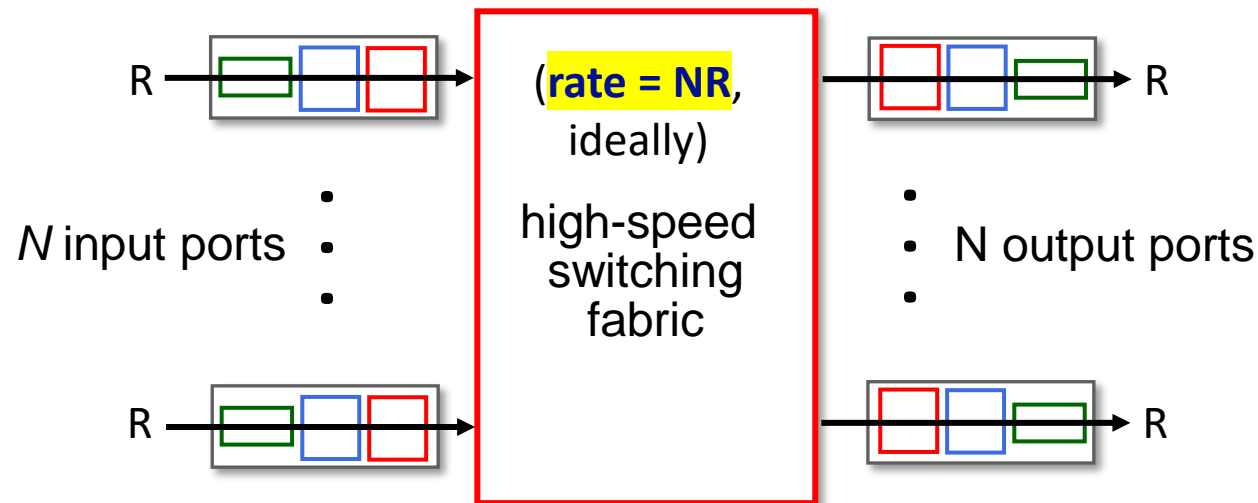
Longest prefix matching



- we'll see *why* longest prefix matching is used shortly, when we study addressing
- longest prefix matching: often performed using ternary content addressable memories (TCAMs)
 - *content addressable*: present address to TCAM: retrieve address in one clock cycle, regardless of table size
 - Cisco Catalyst: ~1M routing table entries in TCAM




Switching fabrics

- transfer packet from input link to appropriate output link
- **switching rate**: rate at which packets can be transfer from inputs to outputs
 - often measured as multiple of input/output line rate
 - For N inputs: switching rate N times line rate R is desirable



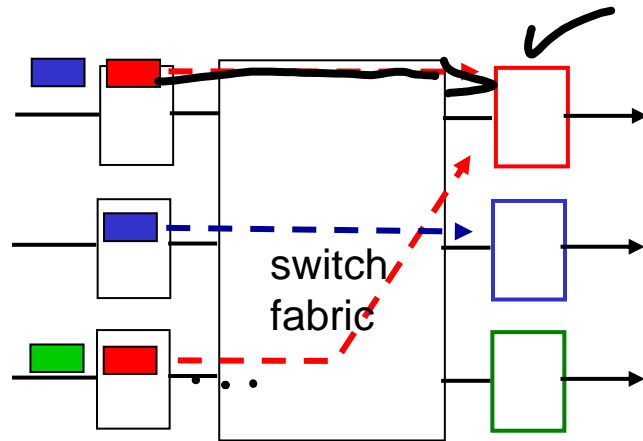
Switching fabrics

- transfer packet from input link to appropriate output link
- **switching rate**: rate at which packets can be transfer from inputs to outputs
 - often measured as multiple of input/output line rate
 - For N inputs: switching rate N times line rate R is desirable
- **three major types of switching fabrics:**




Input port queuing

- If switch fabric slower than input ports combined -> queueing may occur at input queues
 - queueing delay and loss due to input buffer overflow!
- **Head-of-the-Line (HOL) blocking:** queued datagram at front of queue prevents others in queue from moving forward




output port contention: only one red datagram can be transferred. lower red packet is *blocked*



one packet time later: green packet experiences HOL blocking

Output port queuing



This is a really important slide


- **Buffering** required when datagrams arrive from fabric faster than link transmission rate. **Drop policy:** which datagrams to drop if no free buffers?

Datagrams can be lost due to congestion, lack of buffers


- **Scheduling discipline** chooses among queued datagrams for transmission

Priority scheduling – who gets best performance, network neutrality

Output port queuing



at t , packets move
from input to output



one packet time later



- buffering when arrival rate via switch exceeds output line speed
- *queueing (delay) and loss due to output port buffer overflow!*

Queuing

REMEMBER we said that

packets can be “lost within the network” or “dropped at a router.”

- It is *here*, at these *queues* within a router, where such packets are actually dropped and lost.
- The location and extent of queueing (either at the input port queues or the output port queues) will depend on:
 - 1 ■ the traffic load,
 - 2 ■ the speed of the switching fabric, and
 - 3 ■ the line speed.



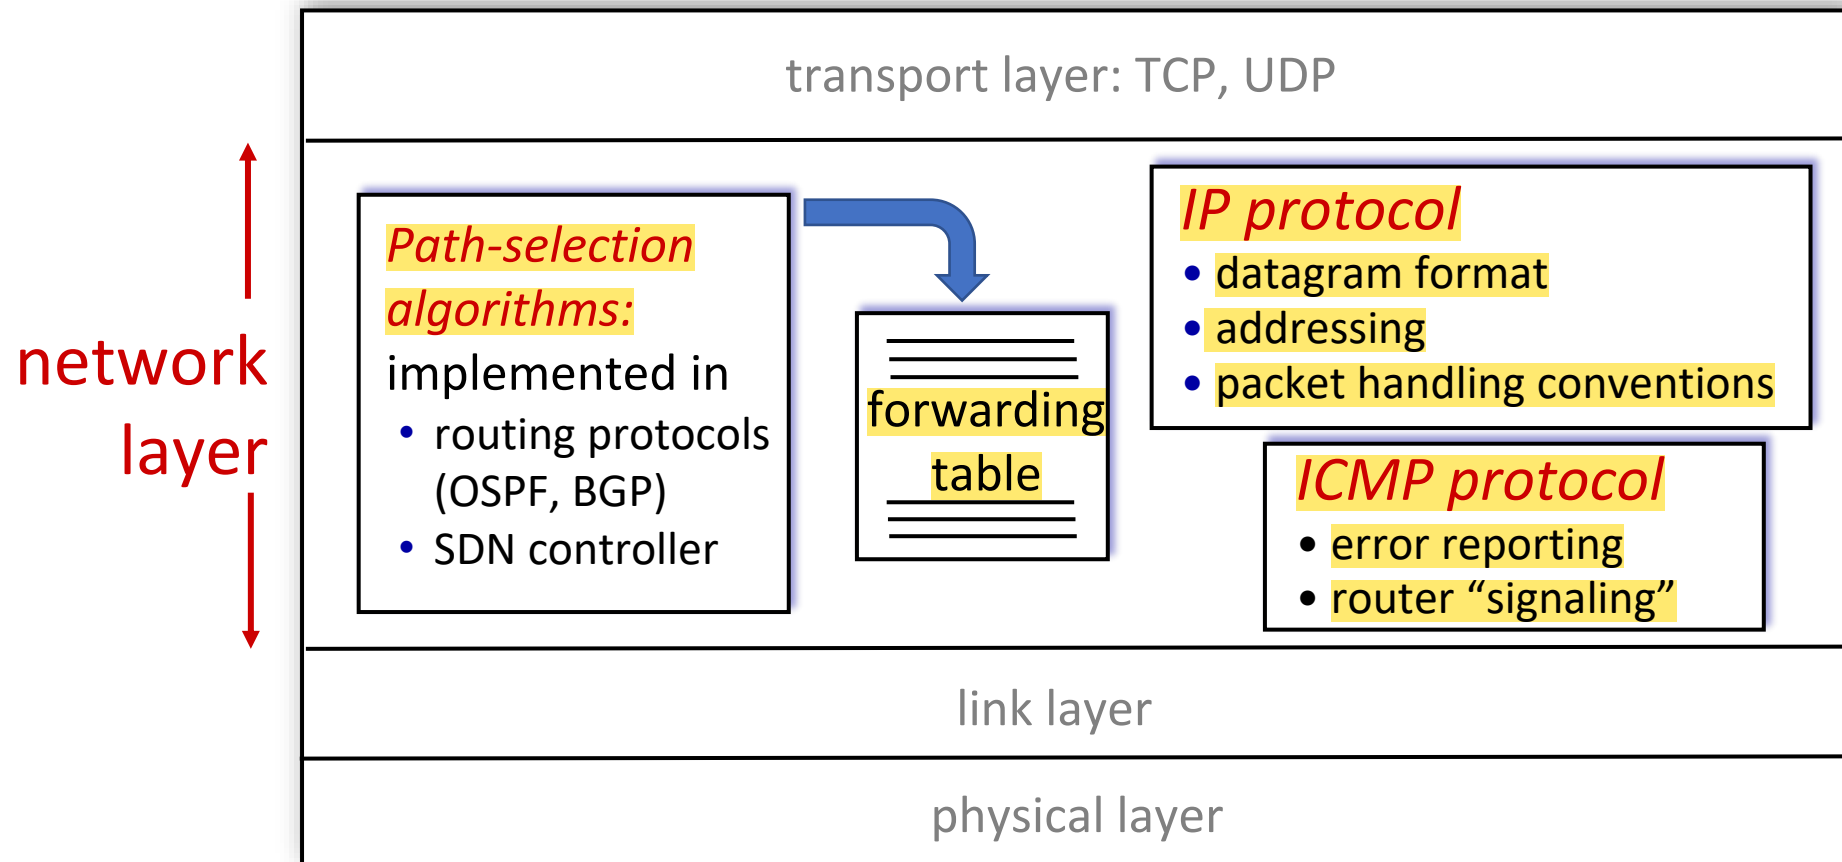
Network layer: “data plane” roadmap

- Network layer: overview
 - data plane
 - control plane
- What’s inside a router
 - input ports, switching, output ports
 - buffer management, scheduling
- IP: the Internet Protocol
 - datagram format
 - addressing
 - network address translation
 - IPv6

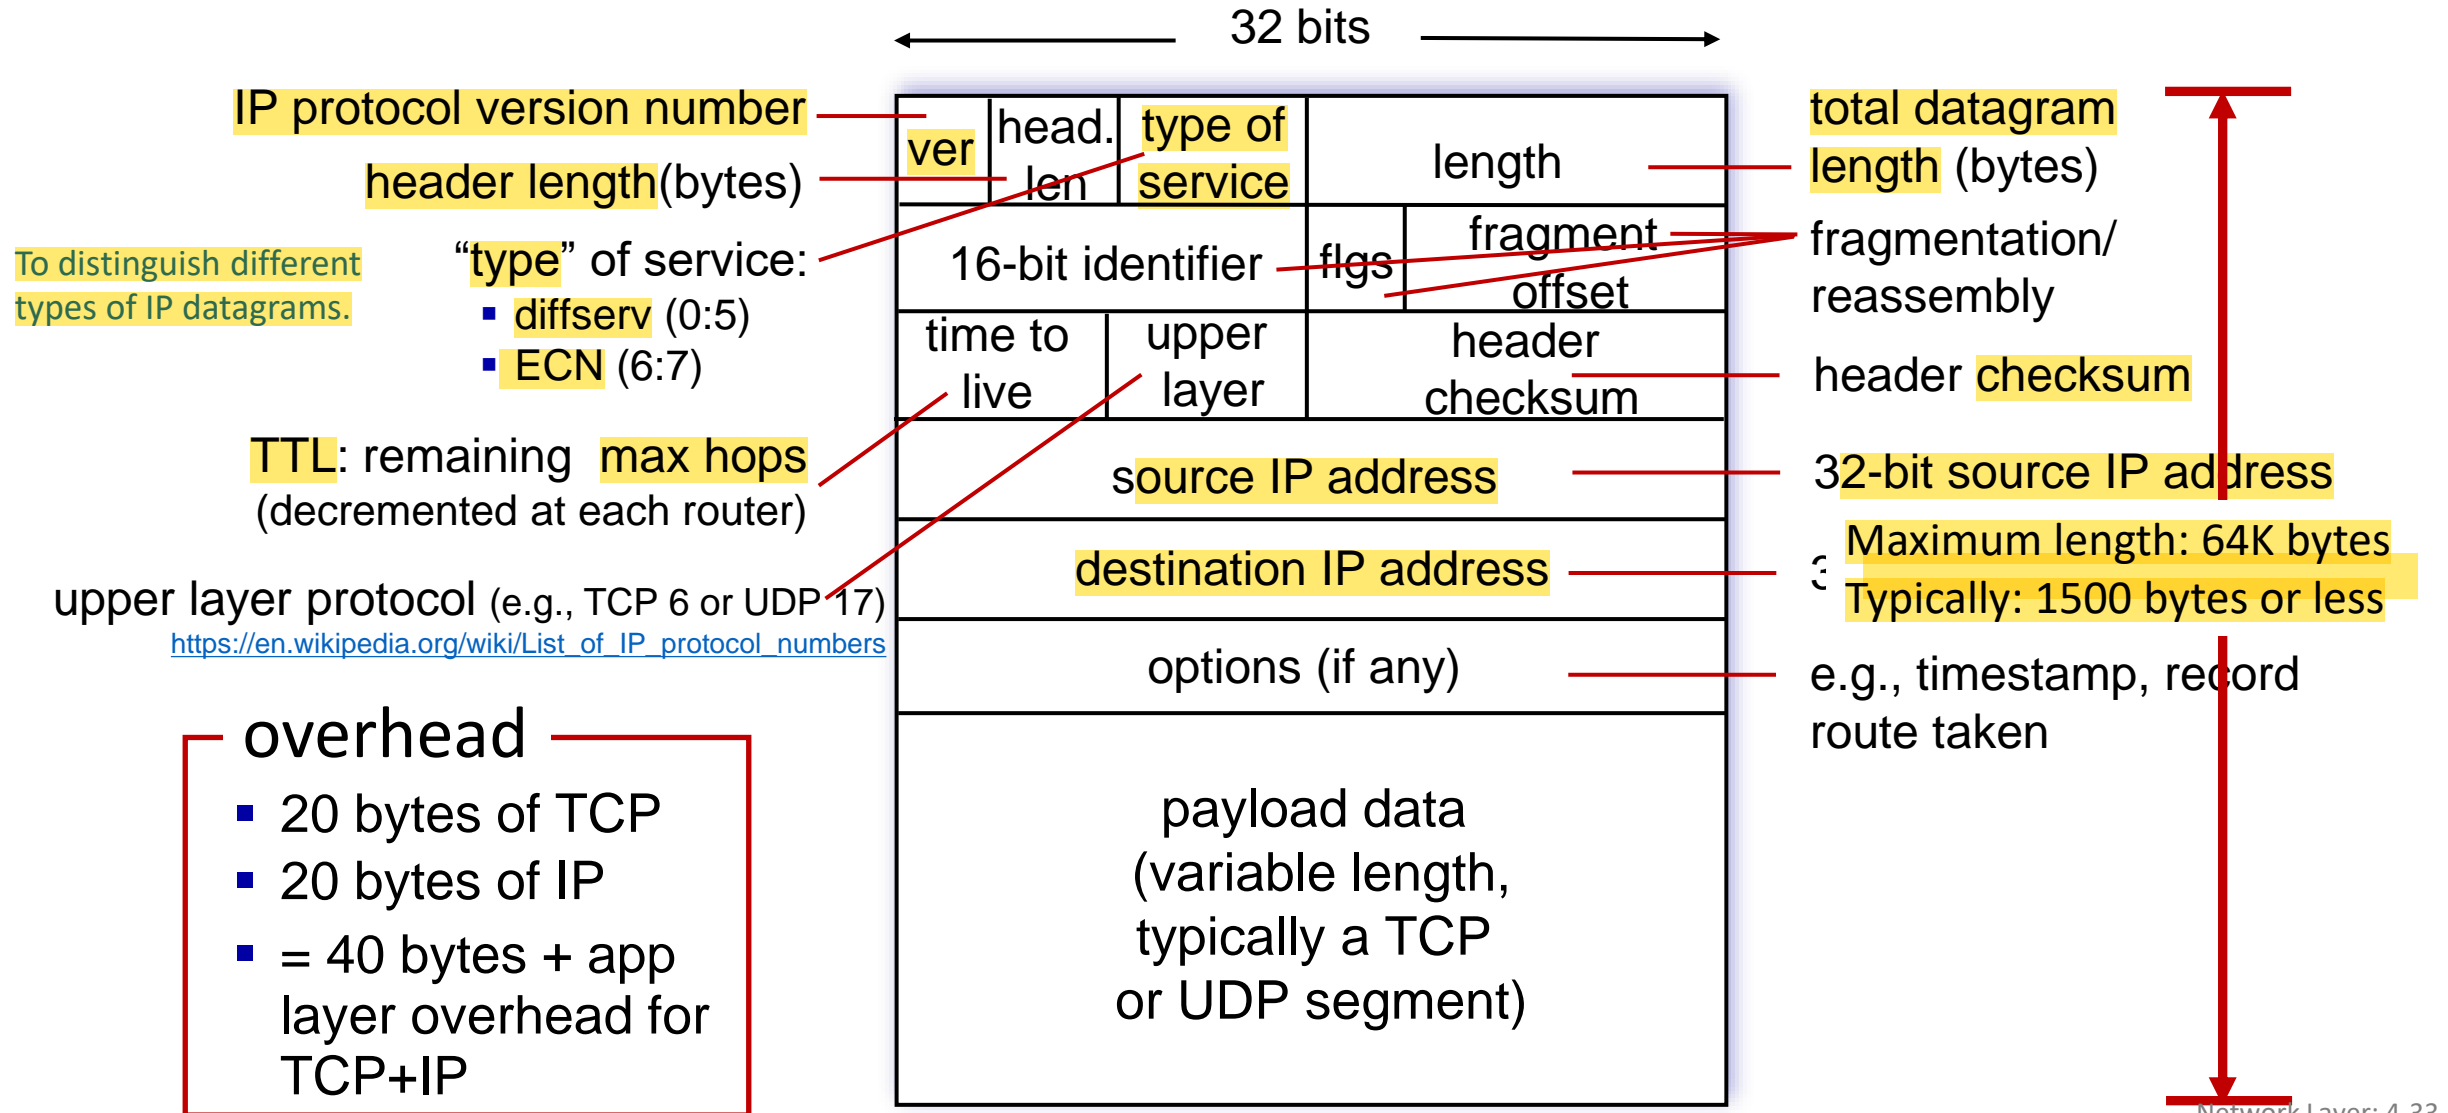


Network Layer: Internet

host, router network layer functions:



IP Datagram format



Network layer: “data plane” roadmap

- Network layer: overview
 - data plane
 - control plane
- What’s inside a router
 - input ports, switching, output ports
 - buffer management, scheduling
- IP: the Internet Protocol
 - datagram format
 - **addressing**
 - network address translation
 - IPv6



IP addressing: introduction

- **IP address:** 32-bit = 4 bytes separated by "." dots (also called 4 octs)
- Each bit can be 0 or 1 (2^1 possibilities)
- 32 bits total possibilities 2^{32} (> 4 billion)
- Each oct has a value between 0 and 255

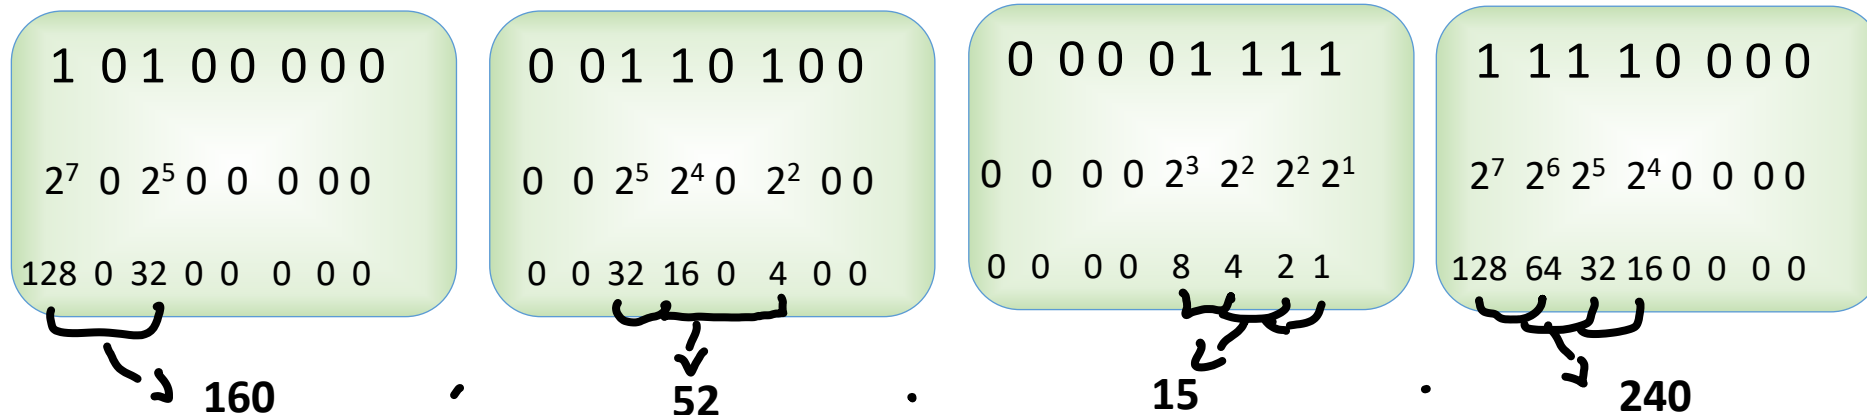
IP address format

8 . 8 . 8 . 8

xxxxxxxx . xxxxxxxx . xxxxxxxx . xxxxxxxx

Ex: 10100000 . 00110100 . 00001111 . 11110000

160 . 52 . 15 . 240



Oct format


X X X X X X X

2^7 2^6 2^5 2^4 2^3 2^2 2^1 2^0

128 64 32 16 8 4 2 1 =255

IP addressing: introduction

- **IP address:** 32-bit identifier associated with each host or router interface
- **interface:** connection between host or router and a physical link
 - router's typically have multiple interfaces
 - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)



dotted-decimal IP address notation:

223.1.1.1 = 11011111 00000001 00000001 00000001

223 1 1 1


IP addressing: introduction

Q: how are interfaces actually connected?

A: we'll learn about that in chapters 6, 7

For now: don't need to worry about how one interface is connected to another (with no intervening router)

A: wired Ethernet interfaces connected by Ethernet switches



A: wireless WiFi interfaces connected by WiFi base station

Subnets


■ What's a subnet ?

- device interfaces that can physically reach each other **without passing through an intervening router**

■ IP addresses have structure:

1. **subnet part:** devices in same subnet have common high order bits
2. **host part:** **remaining** low order bits

A portion of an interface's IP address will be determined by the subnet to which it is connected.



network consisting of 3 subnets

$8 \quad 8 \quad 8 = 24$


223.1.1.1	=	11011111	00000001	00000001	00000001
223.1.1.2	=	11011111	00000001	00000001	00000010
223.1.1.3	=	11011111	00000001	00000001	00000011
223.1.1.4	=	11011111	00000001	00000001	00000100

Sub Network layer: 4-38
host

Subnets

Recipe for defining subnets:

- detach each interface from its host or router, creating “islands” of isolated networks
- each isolated network is called a **subnet**



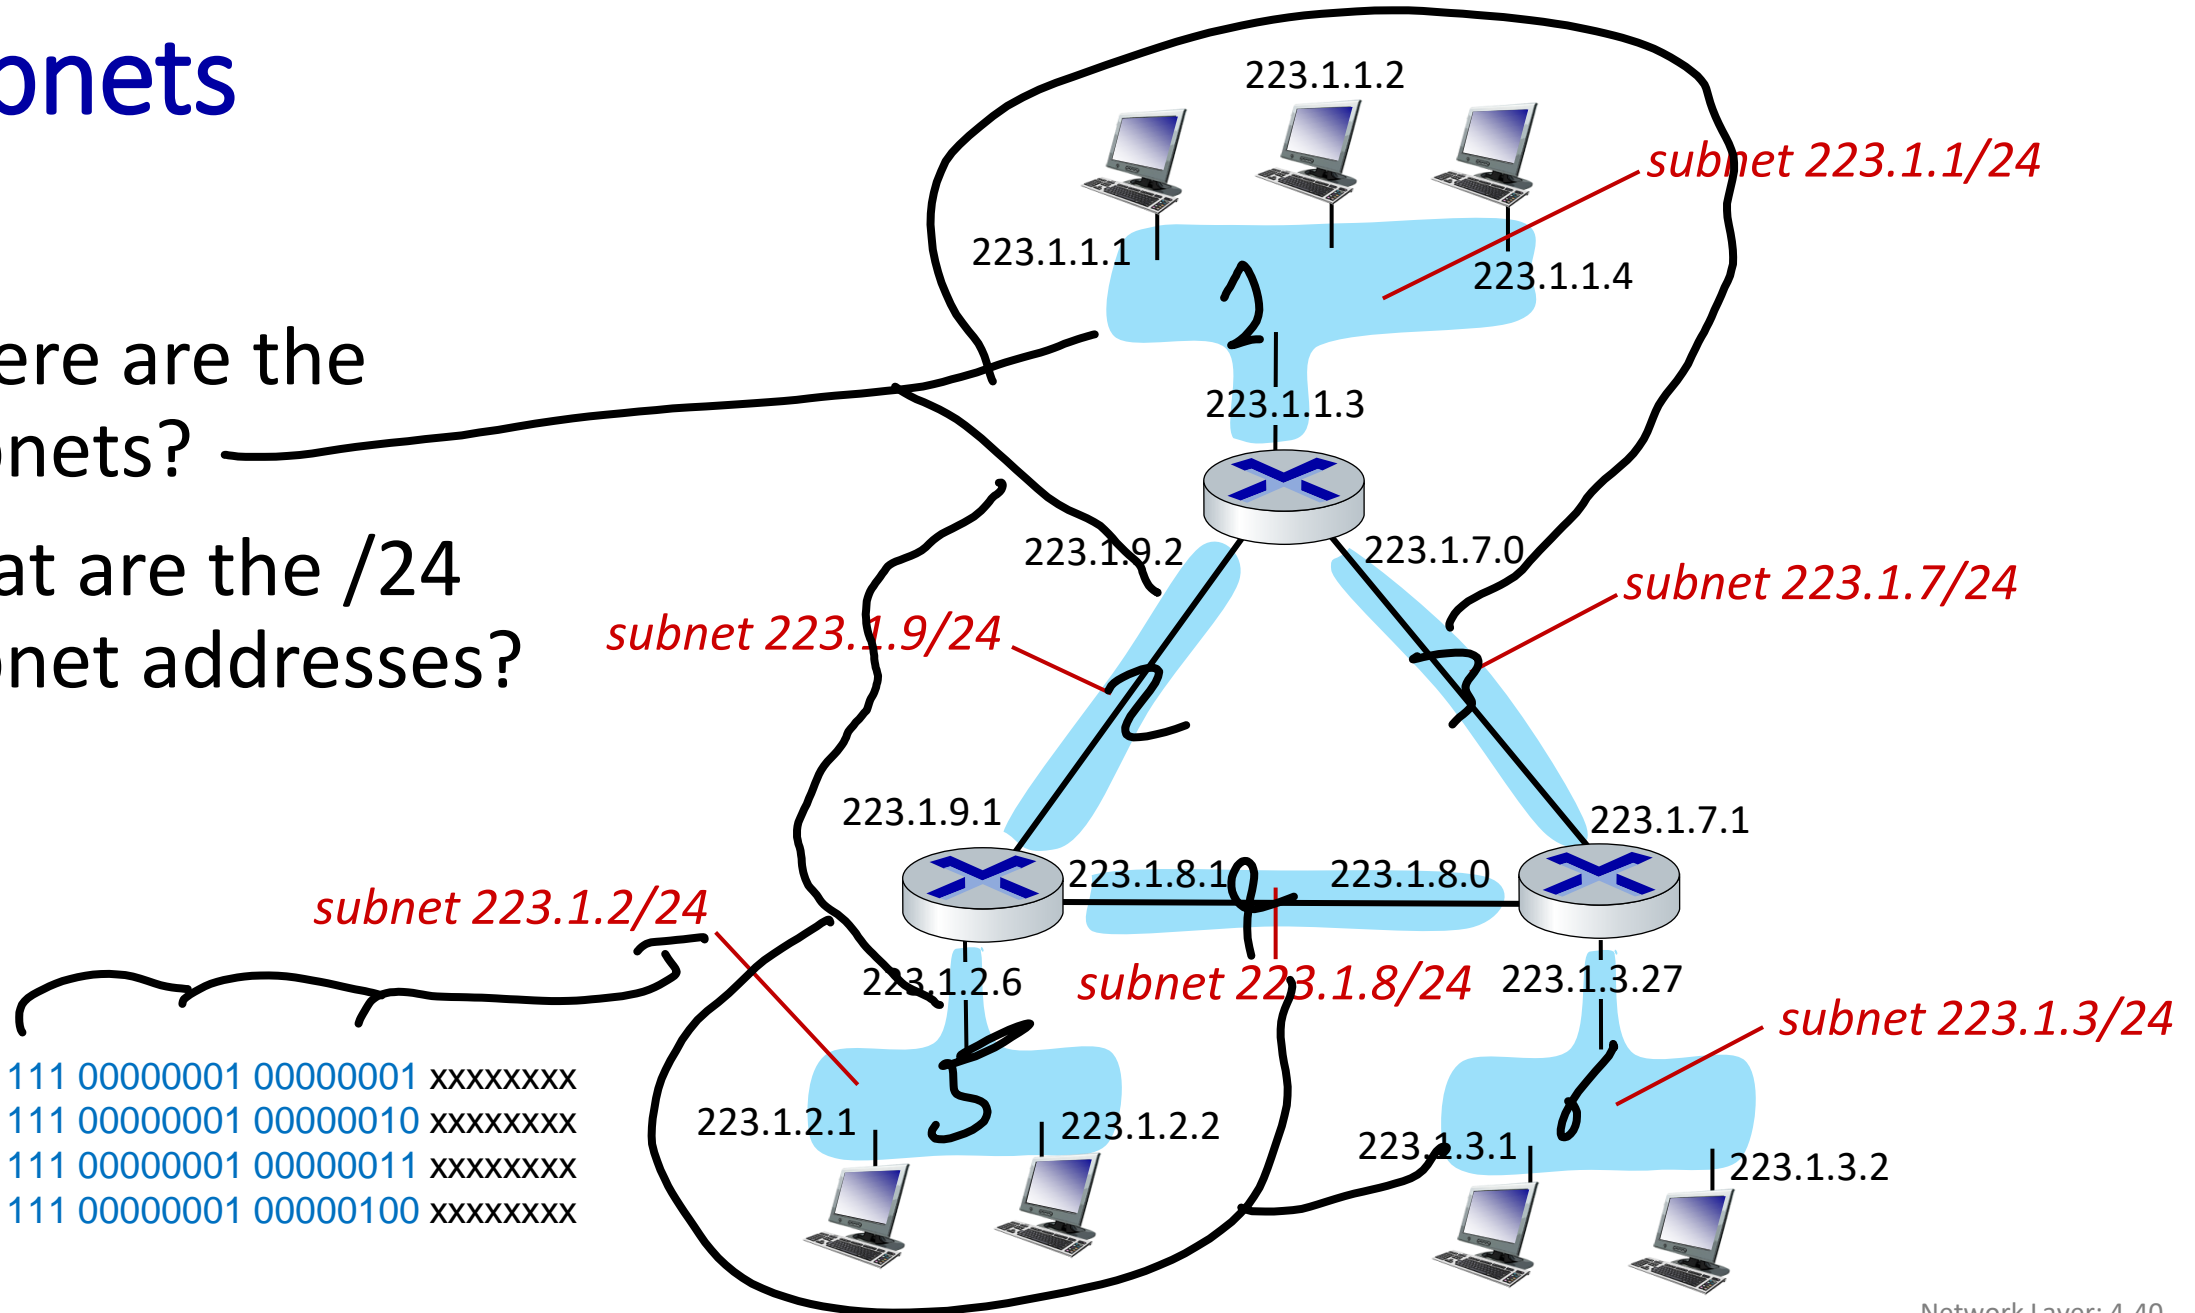
subnet mask: /24

(high-order 24 bits: subnet part of IP address)

Subnets

- where are the subnets?
- what are the /24 subnet addresses?

223.1.1 = 11011111 00000001 00000001 XXXXXXXX
223.1.2 = 11011111 00000001 00000010 XXXXXXXX
223.1.3 = 11011111 00000001 00000011 XXXXXXXX
223.1.4 = 11011111 00000001 00000100 XXXXXXXX



IP Classful addressing

In the past, when an organization requests a range of IP addresses, they receive a from one of these classes:

1 ■ **Class A:** from 0.0.0.1 to 126.0.0.0

- 126 networks and $2^{24} = 16,777,216$ hosts.
- 1 byte for the network & 3 bytes for the host.
- Mask: 255.0.0.0 /8

2 ■ **Class B:** from 128.0.0.0 to 191.255.0.0

- 16,384 networks and $2^{16} = 65,534$ hosts.
- 2 bytes for the network & two for the host.
- Mask: 255.255.0.0 /16

3 ■ **Class C:** from 192.0.0.0 to 223.255.255.0

- 2,097,152 networks and $2^8 = 254$ hosts.
- 3 bytes for the network and the 4th byte for the host.
- Mask: 255.255.255.0 /24

4 ■ **Class D:** from 224.0.0.0 to 239.255.255.
• for multicasting.

5 ■ **Class E:** from 240.0.0.0 to 255.255.255.
• used for experimentation.

Organization requires

20 addresses? C


2000 addresses? B

Which class they get?

IP addressing: CIDR

CIDR: Classless InterDomain Routing (pronounced “cider”)

- subnet portion of address of arbitrary length
- address format: **a.b.c.d/x**, where **x** is # bits in subnet portion of address (**x = mask**)



IP addressing



aaaaaaaa bbbbbbbb cccccccc dddddddd

a.b.c.d/x

Number of IP addresses = 2^{32-x}

a.b.c.d/20	4096	a.b.c.d/26	64 -2
a.b.c.d/21	2048	a.b.c.d/27	32
a.b.c.d/22	1024	a.b.c.d/28	16
a.b.c.d/23	512	a.b.c.d/29	8
a.b.c.d/24	256	a.b.c.d/30	4
a.b.c.d/25	128	a.b.c.d/31	2 -2 = 0 ?

- 1st IP address = Network address
- Last IP address = broadcast address
- Number of IP addresses for hosts = $2^{32-x} - 2$
 - a.b.c.d/x → 4096 IP address & 4094 address for hosts

Example

223 1 1 0


223.1.1.0 = 11011111 00000001 00000001 00000000

128 64 32 16 8 4 2 1

Leftmost bits = 24

223.1.1.0/24

→ **Address Range** 223.1.1.0 to 223.1.1.255



subnet mask: /24

Mask Notation

- Values: *Network = 1* & *Host = 0*
- Classful example (Class B address)
 - 128.35.17.25/ 16 2 8
 - Binary: 11111111 . 11111111 . 00000000 . 00000000
 - Decimal: 255 . 255 . 0 . 0
- Classless IP example
 - 128.35.17.25/ 17 2 2 1
 - Binary: 11111111 . 11111111 . 10000000 . 00000000
 - Decimal: 255 . 255 . 128 . 0

IP addresses: how to get one?

That's actually **two** questions:

1. Q: How does a *host* get IP address within its network (host part of address)?
2. Q: How does a *network* get IP address for itself (network part of address)

How does **host** get IP address?

- hard-coded by sys. admin in config file (e.g., /etc/rc.config in UNIX)
- **DHCP**: Dynamic Host Configuration Protocol: dynamically get address from a server
 - “plug-and-play”

DHCP: Dynamic Host Configuration Protocol


goal: host *dynamically* obtains IP address from network server when it “joins” the network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/on)
- support for mobile users who join/leave network

DHCP overview:

- host broadcasts **DHCP discover** msg [optional]
- DHCP server responds with **DHCP offer** msg [optional]
- host requests IP address: **DHCP request** msg
- DHCP server sends address: **DHCP ack** msg

DHCP client-server scenario



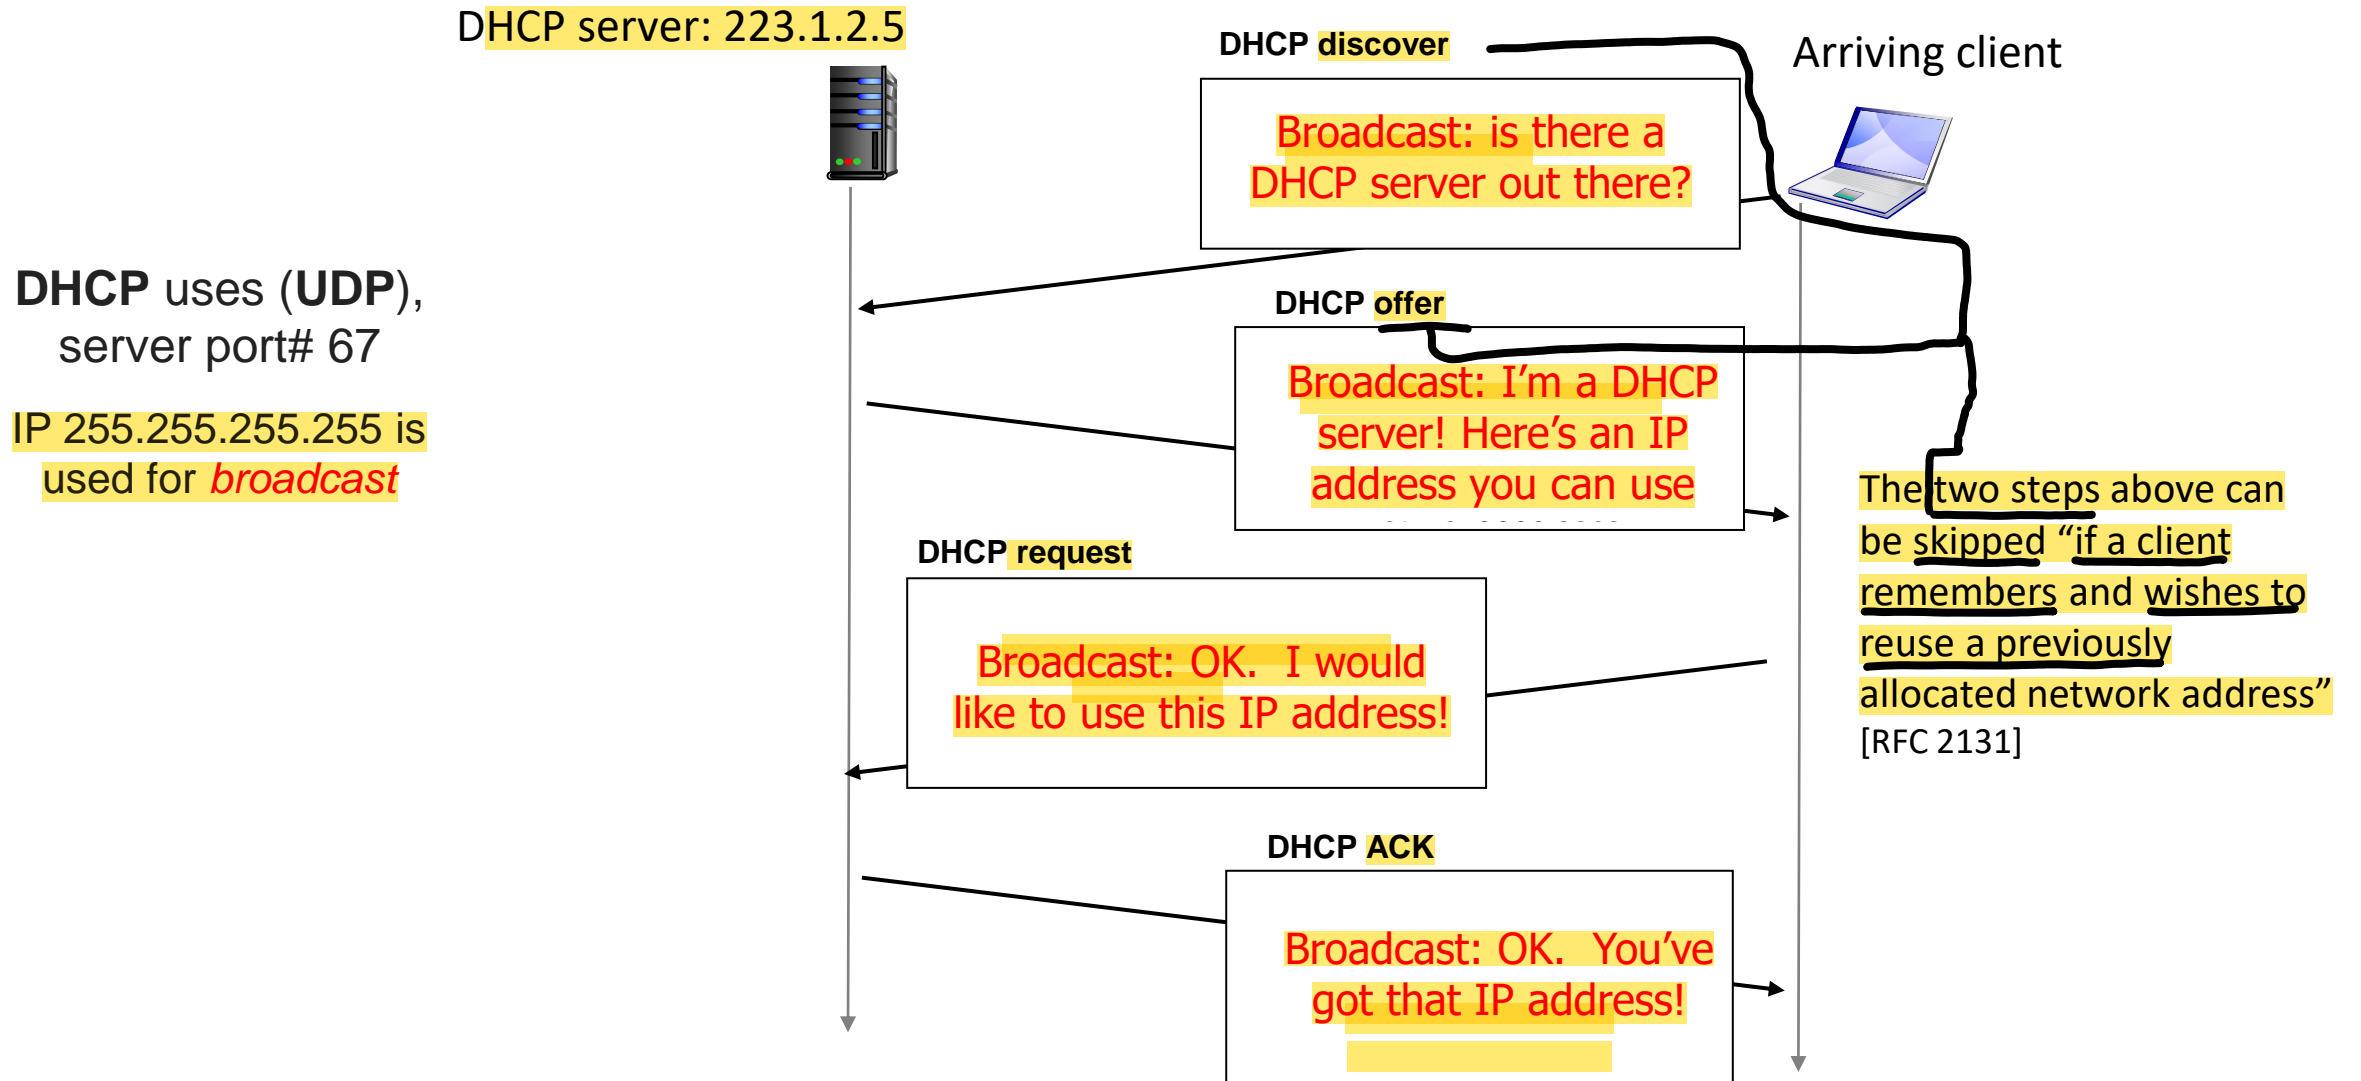
Typically, DHCP server will be co-located in router, serving all subnets to which router is attached



arriving DHCP client needs address in this network

DHCP client-server scenario


Initially, client does not have IP address and does not know server IP address → use 0.0.0.0




DHCP: more than IP addresses

DHCP can return more than just allocated IP address on subnet:

- address of first-hop router for client
- name and IP address of DNS sever
- network mask (indicating network versus host portion of address)




DHCP: example



- Connecting laptop will use DHCP to get IP address, address of first-hop router, address of DNS server.
- DHCP REQUEST message encapsulated in UDP, encapsulated in IP, encapsulated in Ethernet
- Ethernet frame broadcast (dest: FFFFFFFF) on LAN, received at router running DHCP server
- Ethernet demux'ed to IP demux'ed, UDP demux'ed to DHCP

DHCP: example



- DCP server formulates DHCP ACK containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
- encapsulated DHCP server reply forwarded to client, demuxing up to DHCP at client
- client now knows its IP address, name and IP address of DNS server, IP address of its first-hop router

IP addresses: how to get one?

Q: how does *network* get subnet part of IP address?

A: gets allocated portion of its provider ISP's address space

ISP's block $\begin{matrix} 2 & 2 & 9 \end{matrix}$ $\xrightarrow{\hspace{10em}}$ $\begin{matrix} 8+8=16 \\ 16+4=20 \end{matrix}$

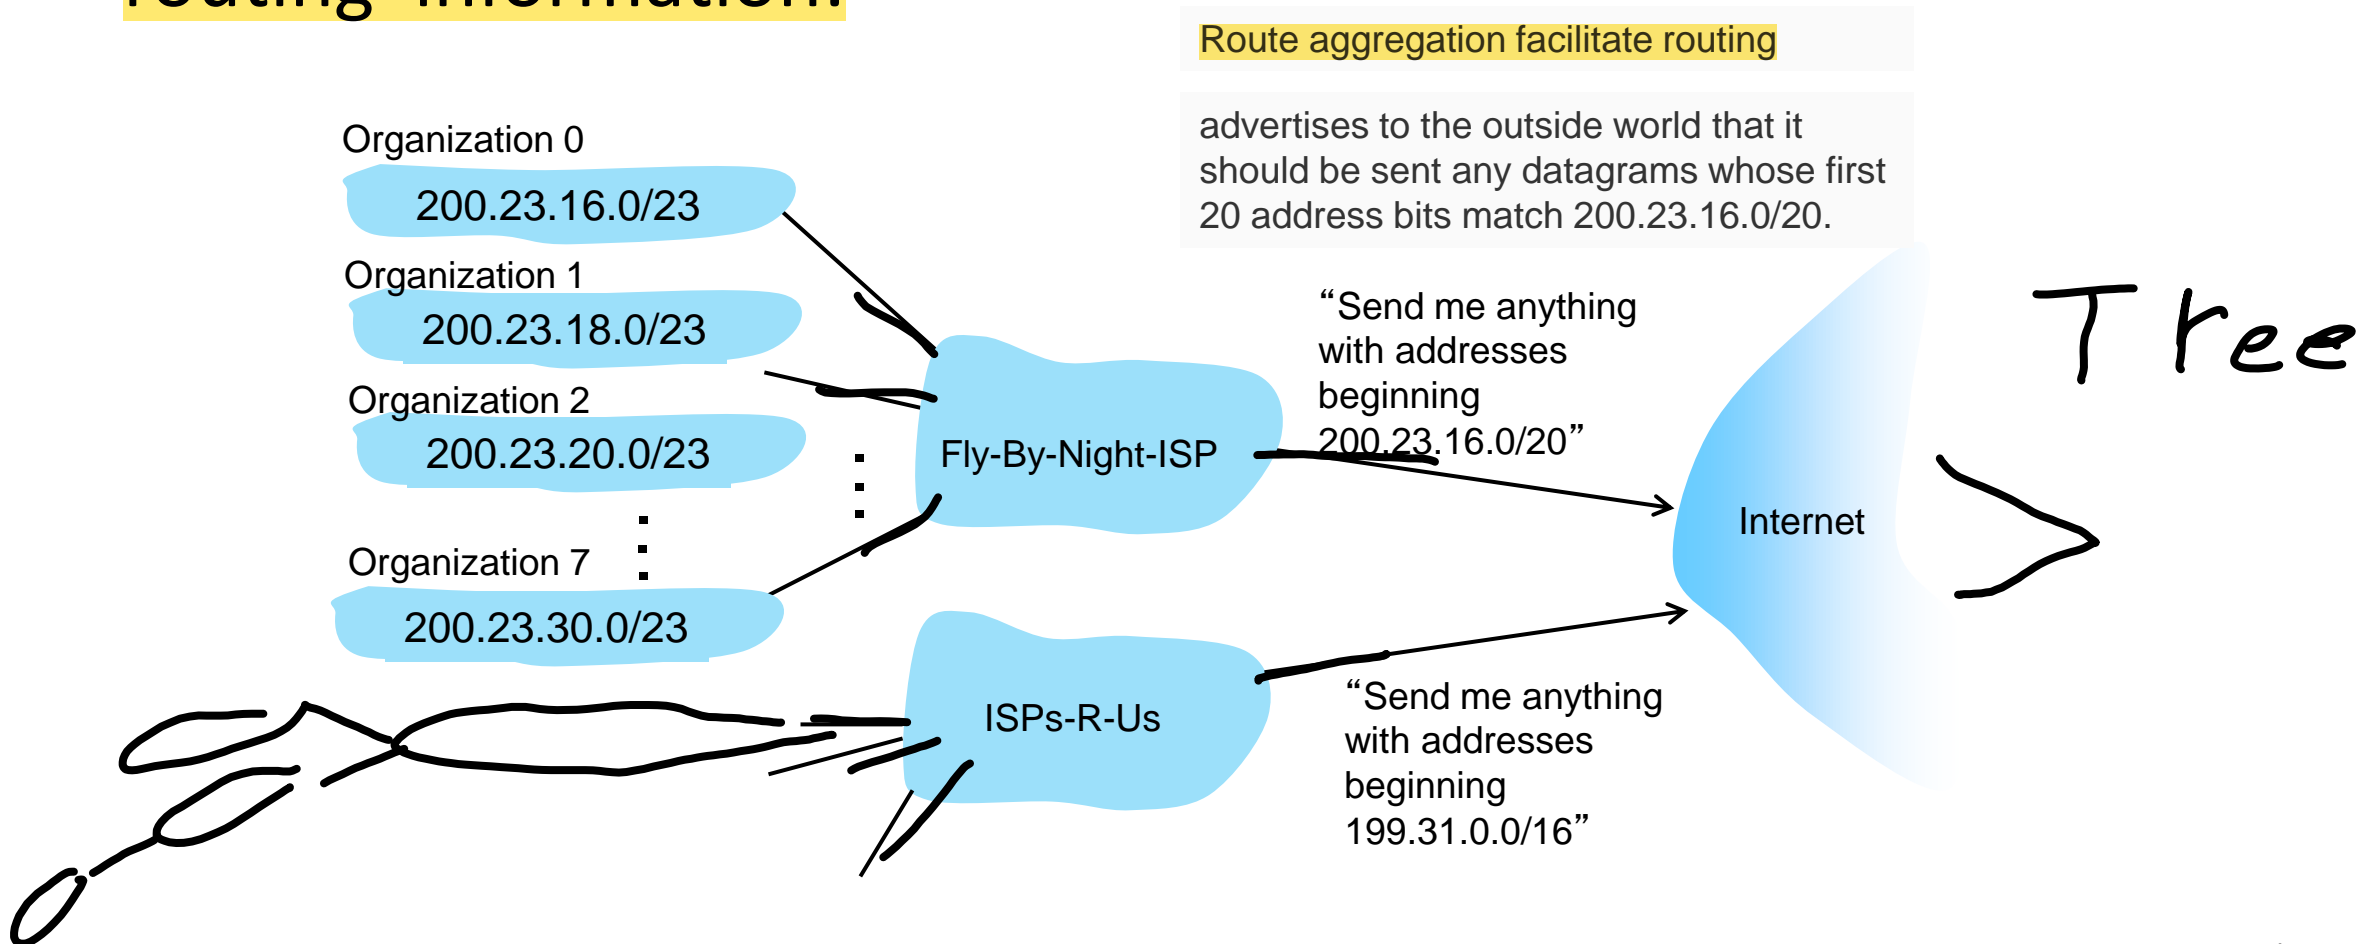
	<u>11001000</u>	<u>00010111</u>	<u>0001</u>	<u>0000</u>	00000000	200.23.16.0/20
--	-----------------	-----------------	-------------	-------------	----------	----------------

ISP can then allocate out its address space in 8 blocks:

	$\begin{matrix} 2 & 2 & 4+3=7 \end{matrix}$ $\xrightarrow{\hspace{10em}}$ $\begin{matrix} 8+8+7=23 \end{matrix}$					
Organization 0	<u>11001000</u>	<u>00010111</u>	<u>0001</u>	<u>000</u>	00000000	200.23.16.0/23
Organization 1	<u>11001000</u>	<u>00010111</u>	<u>0001</u>	<u>001</u>	00000000	200.23.18.0/23
Organization 2	<u>11001000</u>	<u>00010111</u>	<u>0001</u>	<u>010</u>	00000000	200.23.20.0/23
...			
Organization 7	<u>11001000</u>	<u>00010111</u>	<u>0001</u>	<u>111</u>	00000000	200.23.30.0/23

Hierarchical addressing: route aggregation

hierarchical addressing allows efficient advertisement of routing information:



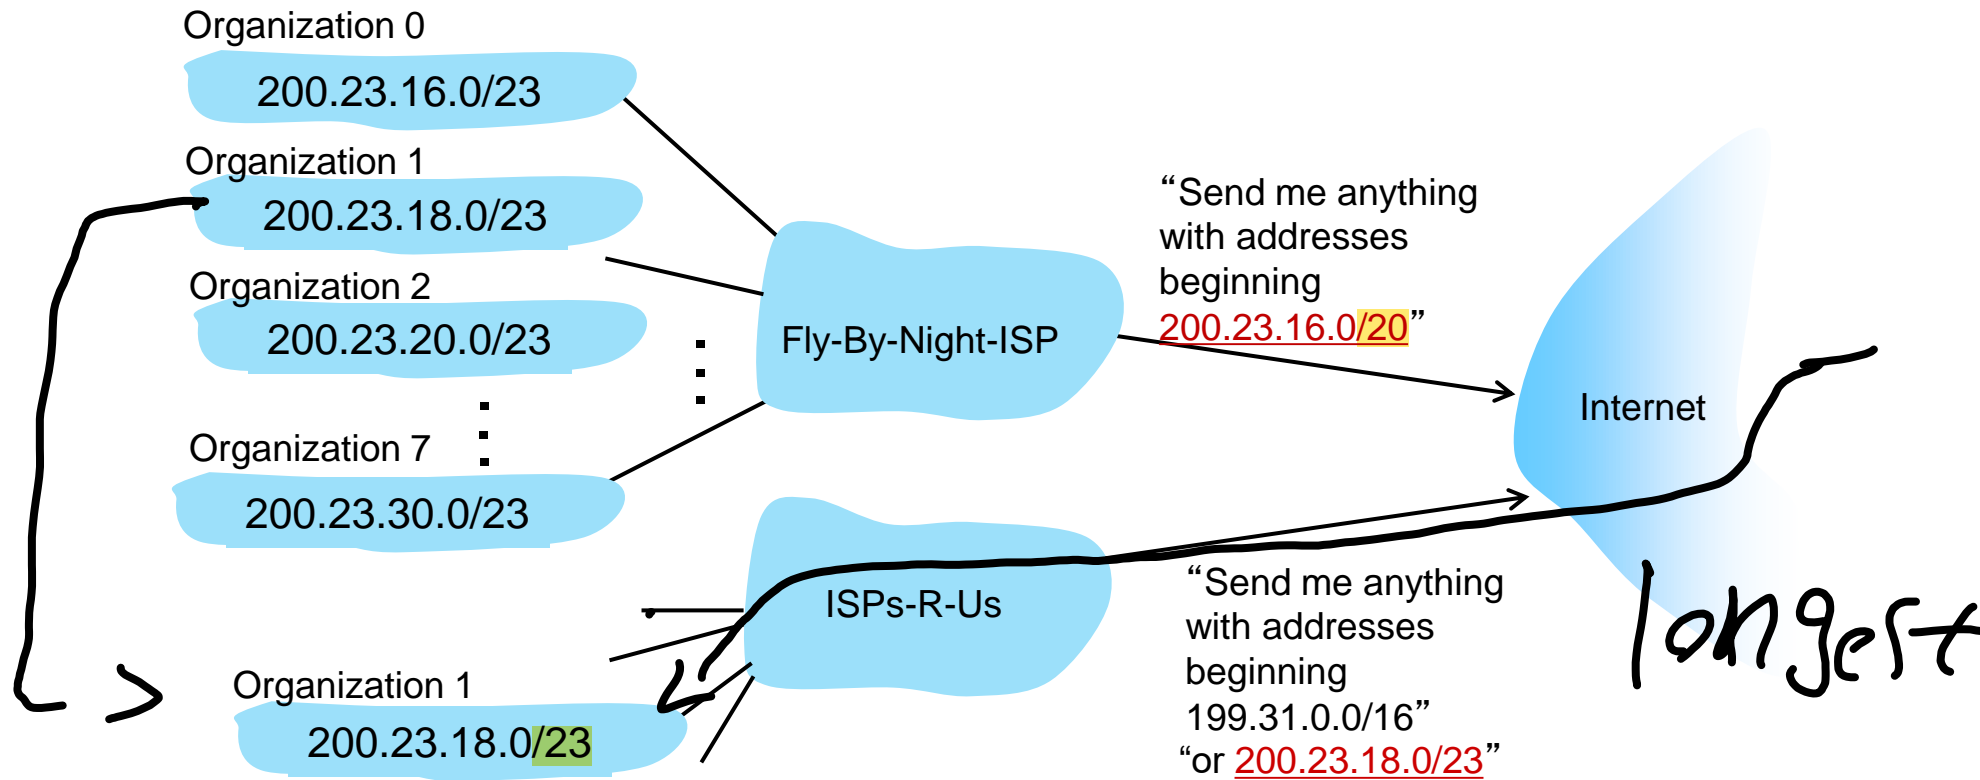
Hierarchical addressing: route aggregation

hierarchical addressing allows efficient advertisement of routing information:

- Hierarchical addressing: addresses are given to ISPs in chunks (continuous set of IP address range) ex. All addresses in a /20 block
 - Then, ISP breaks the big chunk into multiple chunks ex. say /23
 - These chunks can be further broken into smaller chunks ex. /25 and less
 - And so on.
 - All the smaller addresses are part of the original ISP chunk.
- Route aggregation: the ability to use a single prefix to advertise multiple networks is often referred to as address aggregation or route aggregation.

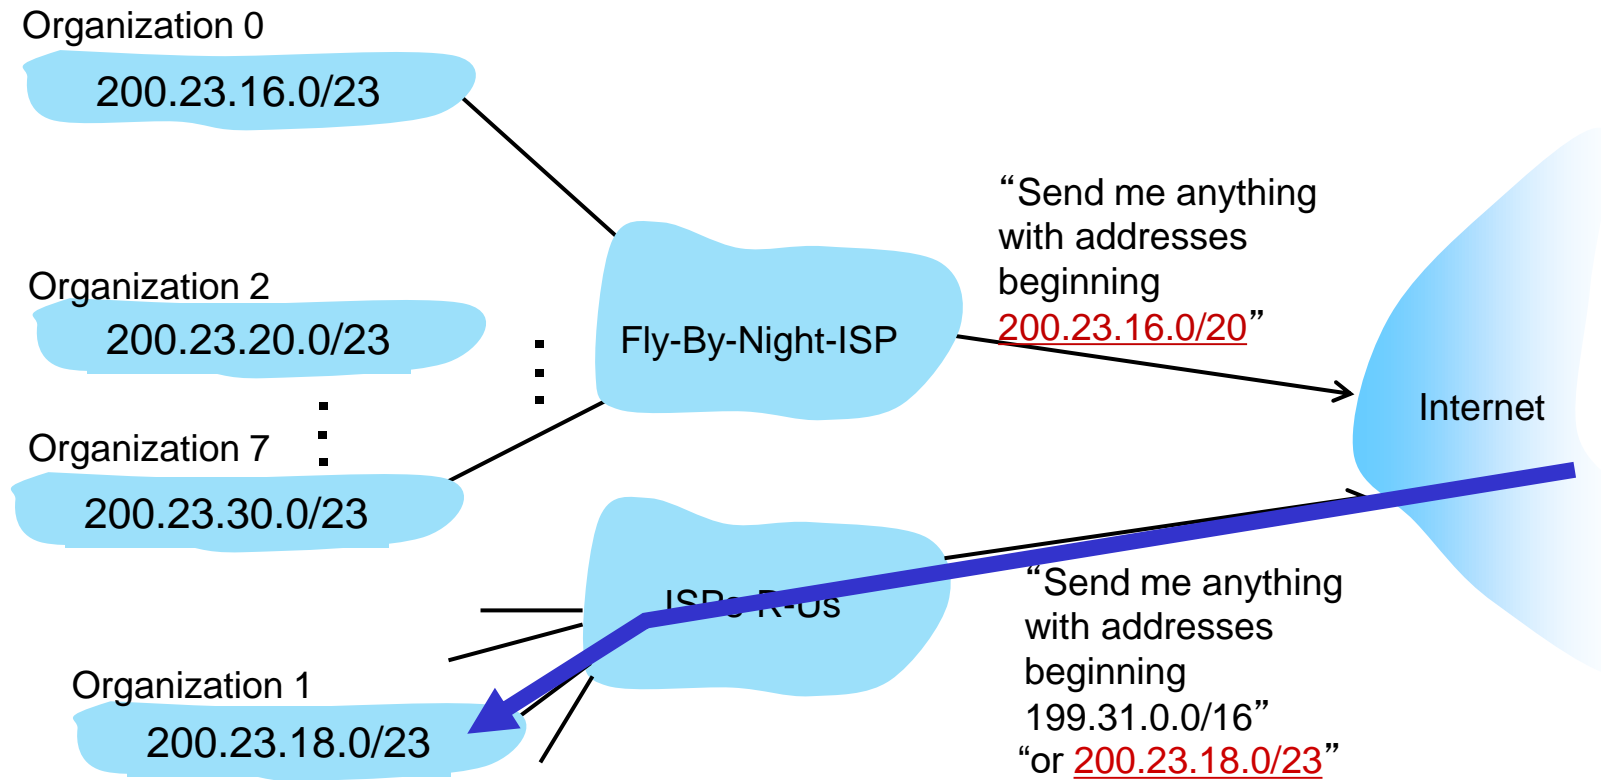
Hierarchical addressing: more specific routes

- If Organization 1 moves from Fly-By-Night-ISP to ISPs-R-Us
- ISPs-R-Us now advertises a more specific route to Organization 1



Hierarchical addressing: more specific routes

- If Organization 1 moves from Fly-By-Night-ISP to ISPs-R-Us
- ISPs-R-Us now advertises a more specific route to Organization 1



IP addressing: last words ...

Q: how does an ISP get block of addresses?

A: ICANN: Internet Corporation for Assigned Names and Numbers
<http://www.icann.org/>

- allocates IP addresses, through 5 regional registries (RRs) (who may then allocate to local registries)
- manages DNS root zone, including delegation of individual TLD (.com, .edu , ...) management

Q: are there enough 32-bit IP addresses?

- There are ($2^{32} > 4.3$) billion address
- ICANN allocated last chunk of IPv4 addresses to RRs in 2011
- NAT (next) helps IPv4 address space exhaustion
- IPv6 has 128-bit address space

"Who the hell knew how much address space we needed?" Vint Cerf (reflecting on decision to make IPv4 address 32 bits long)

Network layer: “data plane” roadmap

- Network layer: overview
 - data plane
 - control plane
- What’s inside a router
 - input ports, switching, output ports
 - buffer management, scheduling
- IP: the Internet Protocol
 - datagram format
 - addressing
 - network address translation
 - IPv6



Private IP addresses

The private IP address ranges follow:


Hosts

- **Class A:** 10.0.0.0 – 10.255.255.255
• 10.0.0.0/8, or just 10/8
 $2^{32-8} \sim 16.8 \text{ million}$
 - **Class B:** 172.16.0.0 – 172.31.255.255
• 172.16.0.0/12, or just 172.16/12
 $2^{32-12} \sim 1.05 \text{ million}$
 - **Class C:** 192.168.0.0 – 192.168.255.255
• 192.168.0.0/16, or just 192.168/16
 $2^{32-16} = 65,536$
- Private IP addresses are not routable outside the local network (they cannot be advertised to the public Internet).
 - They are widely used on almost all local networks today.
 - Private addresses are usually translated with NAT at an edge router to map the private addresses used on a LAN to the public address space used by the ISP.

Local Network (Home, Company, or university, etc)

Private IP addresses


10/8,
172.16/12,
or 192.168/16



- Private IP addresses are not routable outside the local network (they cannot be advertised to the public Internet).
- They are widely used on almost all local networks today.
- Private addresses are usually translated with NAT at an edge router to map the private addresses used on a LAN to the public address space used by the ISP.

NAT: network address translation

NAT: all devices in local network share just **one** IPv4 address as far as outside world is concerned



all datagrams *leaving* local network have *same* source NAT IP address: 138.76.29.7, but *different* source port numbers

datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)

NAT: network address translation

- all devices in local network have 32-bit addresses in a “private” IP address space (10/8, 172.16/12, 192.168/16 prefixes) that can only be used in local network
- advantages:
 - 1 ■ just **one** IP address needed from provider ISP for *all* devices
 - 2 ■ can change addresses of host in local network without notifying outside world
 - 3 ■ can change ISP without changing addresses of devices in local network
 - 4 ■ security: devices inside local net not directly addressable, visible by outside world

NAT: network address translation

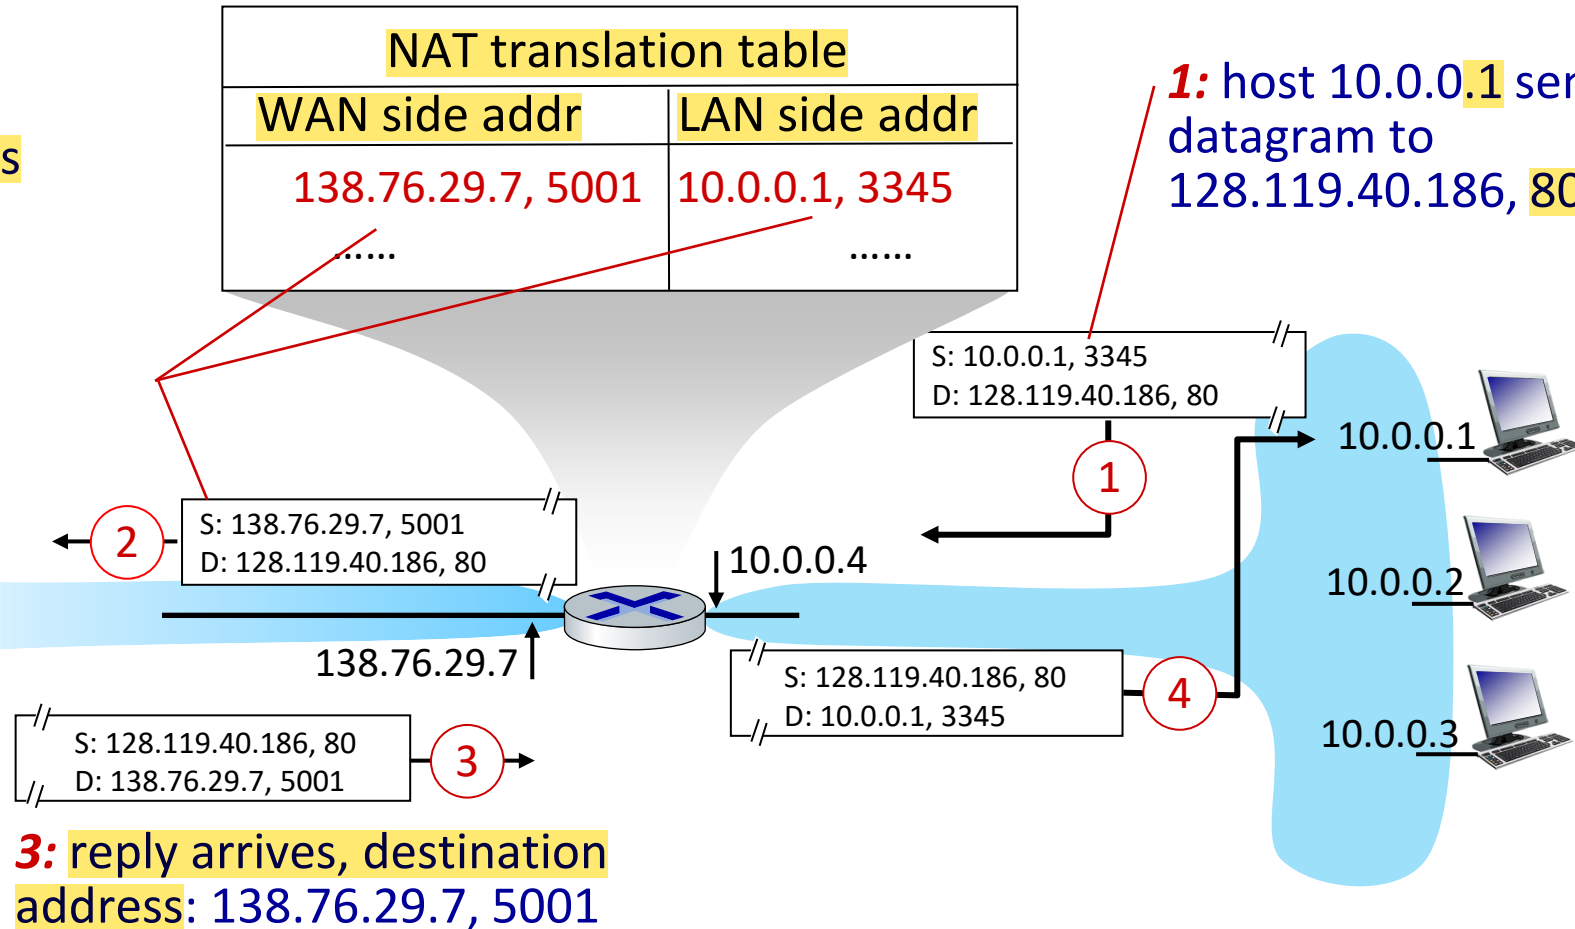
implementation: NAT router must (transparently):

- outgoing datagrams: replace (source IP address, socket port #) of every outgoing datagram to (NAT IP address, new port #)
 - remote clients/servers will respond using (NAT IP address, new port #) as destination address
- remember (in NAT translation table) every (source IP address, port #) to (NAT IP address, new port #) translation pair
- incoming datagrams: replace (NAT IP address, new port #) in destination fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

NAT: network address translation

2: NAT router changes datagram source address from 10.0.0.1, 3345 to 138.76.29.7, 5001, updates table

1: host 10.0.0.1 sends datagram to 128.119.40.186, 80



NAT: network address translation

- NAT has been controversial:
 - routers “should” only process up to layer 3 socket port # → transport layer
 - address “shortage” should be solved by IPv6 2^{128} → 3.4×10^{38}
 - violates end-to-end argument (port # manipulation by network-layer device)
 - NAT traversal: what if client wants to connect to server behind NAT?
- but NAT is here to stay:
 - extensively used in home and institutional nets, 4G/5G cellular nets

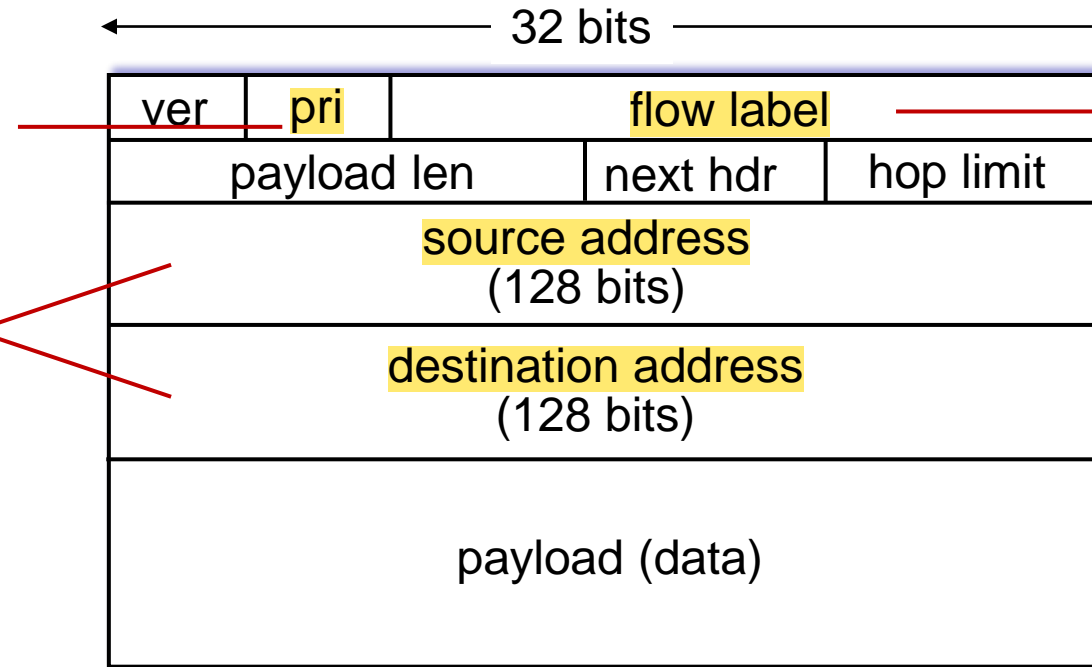
IPv6: motivation

- **initial motivation:** 32-bit IPv4 address space would be completely allocated
- **additional motivation:**
 - 1 • speed processing/forwarding: 40-byte fixed length header
 - 2 • Enable different network-layer treatment of “flows”

IPv6 datagram format

priority: identify priority among datagrams in flow

128-bit IPv6 addresses



flow label: identify datagrams in same "flow." (concept of "flow" not well defined).

What's missing (compared with IPv4):


- 1 ■ no checksum (to speed processing at routers)
- 2 ■ no fragmentation/reassembly
- 3 ■ no options (available as upper-layer, next-header protocol at router)

IPv6: adoption

- Google¹: ~ 30% of clients access services via IPv6
- NIST: 1/3 of all US government domains are IPv6 capable

IPv6 Adoption

We are continuously measuring the availability of IPv6 connectivity among Google users. The graph shows the percentage of users that access Google over IPv6.



1

<https://www.google.com/intl/en/ipv6/statistics.html>

IPv6: adoption

- Google¹: ~ 30% of clients access services via IPv6
- NIST: 1/3 of all US government domains are IPv6 capable
- Long (long!) time for deployment, use
 - 25 years and counting!
 - think of application-level changes in last 25 years: WWW, social media, streaming media, gaming, telepresence, ...
 - *Why?*

¹ <https://www.google.com/intl/en/ipv6/statistics.html>

A black mug of coffee sits on a wooden table. Steam rises from the mug. Coffee beans are scattered on the table surface. The background is a warm, brownish-orange wall.


Revision

Additional IP addressing problems

Additional Problems

IP addressing

IP addressing



a.b.c.d/x Number of hosts = 2^{32-x}

a.b.c.d/20	4096	a.b.c.d/26	64
a.b.c.d/21	2048	a.b.c.d/27	32
a.b.c.d/22	1024	a.b.c.d/28	16
a.b.c.d/23	512	a.b.c.d/29	8
a.b.c.d/24	256	a.b.c.d/30	4
a.b.c.d/25	128	a.b.c.d/31	2

Oct format

X X X X X X X

$2^7 \ 2^6 \ 2^5 \ 2^4 \ 2^3 \ 2^2 \ 2^1 \ 2^0$

128 64 32 16 8 4 2 1 =255

Usually, first and last IP addresses are reserved for network address (or subnet) and broadcast address

IP addressing

- **Network address** - the address by which we refer to the network.
 - the first IP address in it which all host part bits = 0
 - Network Addresses have all 0's in the host portion.
- **Broadcast address** - A special address used to send data to all hosts in the network
 - the last IP address in the network which all host part bits = 1
 - Broadcast Addresses have all 1's in the host portion.
- **Host addresses** - The addresses assigned to the end devices in the network
 - Host Addresses can not have all 0's or all 1's in the host portion.

example:

223.1.1.0/24

11011111 00000001 00000001 00000000

Number of IP addresses: $2^{32-24} = 2^8 = 256$ IP address

Network address = 223.1.1.0


broadcast = 223.1.1.255

Hosts take the remaining addresses from 223.1.1.1 to 223.1.1.254 → $256-2 = 254$ hosts

IP addressing

- Which IP address should be assigned to PC B ?


- a) 192.168.15.5 X 3^{ed} byte "15"
- b) 192.168.5.0 X subnet address
- c) 192.168.5.40 ✓
- d) 192.168.5.255 X broadcast
- e) 192.168.5.256 X error
?



IP addressing

- Which IP address should be assigned to PC B ?

- a) 172.16.255.255 X broadcast
- ☒ b) 172.16.0.255 ✓
- c) 172.16820.21 X another subnet
- ☒ d) 172.16.254.255 ✓
- e) 172.16.0.0 X subnet address



IP addressing

- Which IP address should be assigned to PC B ?

a) 192.168.5.5 ✗

b) 192.168.5.0 ✗ Sub

c) 192.168.5.66 ✗

d) 192.168.5.255 ✗

e) 192.168.5.50

We can't tell !!

Need to convert to binary

192. 168. 5. 60 /28

11000000 . 10101000 . 00000101 . 00111100

Network

11000000 . 10101000 . 00000101 . 0011 0000


192.168.5.48/28

Broadcast

11000000 . 10101000 . 00000101 . 0011 1111

192.168.5.63/28

Host addresses = any IP between these: 192.168.5.48 to 192.168.5.63



IP addressing

- Which IP address should be assigned to PC B ?

a) 192.168.5.5 ✗

b) 192.168.5.0 *sub*

c) 192.168.5.62

d) 192.168.5.255

e) 192.168.5.50 ✗

We can't tell !!

Need to convert to binary

192. 168. 5. 61 /30

11000000 . 10101000 . 00000101 . 00111101

Network

11000000 . 10101000 . 00000101 . 00111100

192.168.5.60/30


Broadcast

11000000 . 10101000 . 00000101 . 00111111

192.168.5.63/30

Host addresses = any IP between these: 192.168.5.60 to 192.168.5.63

A: 192.168.5.61/30



B: ???

IP addresses: how to get one?

Q: how does *network* get subnet part of IP address?

A: gets allocated portion of its provider ISP's address space

ISP's block 11001000 00010111 00010000 00000000 200.23.16.0/20

ISP can then allocate out its address space in 8 blocks:

Organization 0	<u>11001000 00010111 0001</u> 0000	00000000	200.23.16.0/23
Organization 1	<u>11001000 00010111 0001</u> 0010	00000000	200.23.18.0/23
Organization 2	<u>11001000 00010111 0001</u> 0100	00000000	200.23.20.0/23
...
Organization 7	<u>11001000 00010111 0001</u> 1110	00000000	200.23.30.0/23

IP addresses: how to get one?

ISP's block 11001000 00010111 00010000 00000000 200.23.16.0/20

What if ISP want to divide its address space into 4 blocks:

Organization 0	<u>11001000 00010111 0001</u> 0000 00000000	200.23.16.0/22
Organization 1	<u>11001000 00010111 0001</u> 0100 00000000	200.23.20.0/22
Organization 2	<u>11001000 00010111 0001</u> 1000 00000000	200.23.24.0/22
Organization 3	<u>11001000 00010111 0001</u> 1100 00000000	200.23.28.0/22

IP addresses: how to get one?

ISP's block 11001000 00010111 00010000 00000000 200.23.16.0/20

What if ISP want to divide its address space into 16 blocks:

Organization 0	<u>11001000 00010111 0001</u>	0000	00000000	200.23.16.0/24
Organization 1	<u>11001000 00010111 0001</u>	0001	00000000	200.23.17.0/24
Organization 2	<u>11001000 00010111 0001</u>	0010	00000000	200.23.18.0/24
...
...
Organization 16	<u>11001000 00010111 0001</u>	1111	00000000	200.23.31.0/24

Consider a router that interconnects three subnets: Subnet 1, Subnet 2, and Subnet 3.

Suppose all the three subnets are required to have the prefix 223.1.17/24.

Suppose that Subnet 1 is required to support at least 60 interfaces,

Subnet 2 is to support at least 90 interfaces, and

Subnet 3 is to support at least 12 interfaces.

Provide three network addresses (of the form a.b.c.d/x)

that satisfy these constraints.

223.1.17/24

→

256

Subnet 1 > 60 interfaces

→

Closest is 64 → mask /26

Subnet 2 > 90 interfaces

→

Closest is 128 → mask /25

Subnet 3 > 12 interfaces

→

Closest is 16 → mask /28

<u>/mask</u>	<u># host</u>
a.b.c.d/20	4096
a.b.c.d/21	2048
a.b.c.d/22	1024
a.b.c.d/23	512
a.b.c.d/24	256
a.b.c.d/25	128
a.b.c.d/26	64
a.b.c.d/27	32
a.b.c.d/28	16
a.b.c.d/29	8
a.b.c.d/30	4
a.b.c.d/31	2

Continue....

Oct format

223.1.17/24

11011111 00000001 00010001 xxxxxxxx

X	X	X	X	X	X	X	X
2 ⁷	2 ⁶	2 ⁵	2 ⁴	2 ³	2 ²	2 ¹	2 ⁰
128	64	32	16	8	4	2	1

Subnet 1 (/26)

11011111 00000001 00010001 00 xxxxxx

223.1.17.0 /26

00 000000

223.1.17.0 566

00 111111

223.1.17.63 640

Subnet 2 (/25)

11011111 00000001 00010001 1 xxxxxxxx

223.1.17.128 /25

1 0000000

223.1.17.128 566

1 1111111

223.1.17.255 640

Subnet 3 (/28)

11011111 00000001 00010001 0100 xxxx

223.1.17.64 /28

or, 0110 xxxx

223.1.17.96 /28

or, 0111 xxxx

223.1.17.112 /28

Subnetting

- Use this block: 64.1.1/24
- Each subnet should have 32 IP addresses
- Subnets between routers → 2 IP addresses
 - $2^{32-27} = 2^5 = 32$ host addresses
 - $2^{32-31} = 2$ host addresses

01000000 00000001 00000001 xxxxxxxx

01000000 00000001 00000001 000xxxxx
01000000 00000001 00000001 001xxxxx
01000000 00000001 00000001 010xxxxx

01000000 00000001 00000001 0110000x
01000000 00000001 00000001 0110001x
01000000 00000001 00000001 0110010x


64.1.1/24

64.1.1.0/27
64.1.1.32/27
64.1.1.64/27

64.1.1.96/31
64.1.1.98/31
64.1.1.100/31

subnet 64.1.1.32/27

subnet 64.1.1.64/27



Forwarding decision

Q: Given this forwarding table, computer the forwarding decision for the following:

16.16.20.28 00010000 00010000 00010100 00011100

16.16.4.28 00010000 00010000 0000 0100 00011100

16.100.100.234 00010000 01100100 01100100 11101010

1

2

5

Forwarding Table	
Dest. Address range	Output interface
16.16.0.0/16	1
16.16.0.0/20	2
64.1.0.0/20	3
64.1.1.0/24	4
16 /8	5

Preparation for answer			1 st step: convert to binary (easy way)	
128 64 32 16 8 4 2 1	128 64 32 16 8 4 2 1	128 64 32 16 8 4 2 1		
<u>00010000</u>	<u>00010000</u>			1
<u>00010000</u>	<u>00010000</u>	<u>0000</u> xxxx		2
01000000	0000000 <u>1</u>	0000 xxxx		3
01000000	00000001	0000 0001		4
00010000				5

Alternative methods to subnetting

If you like it, you can use it.

Alternative approach to Subnetting

Borrow 1 bit from the host portion of the address.

Original	192.	168.	1.	0	000	0000
Mask	255.	255.	255.	0	000	0000
						1 Network

The borrowed bit value is 0 for the Net 0 address.

Net 0	192.	168.	1.	0	000	0000	
Net 1	192.	168.	1.	1	000	0000	
							2 Subnets

The new subnets have the **SAME** subnet mask.

Mask	255.	255.	255.	1	000	0000
------	------	------	------	---	-----	------

Decimal Representation

Original	192.	168.	1.	0	000	0000	Network: 192.168.1.0/24
Mask	255.	255.	255.	0	000	0000	Mask: 255.255.255.0

Borrowing 1 bit creates 2 subnets with the same mask.

Net 0	192.	168.	1.	0	000	0000	Network: 192.168.1.0/25
Mask	255.	255.	255.	1	000	0000	Mask: 255.255.255.128
Net 1	192.	168.	1.	1	000	0000	Network: 192.168.1.128/25
Mask	255.	255.	255.	1	000	0000	Mask: 255.255.255.128

IP addressing

- The prefix and the subnet mask are different ways of representing the same thing - the network portion of an address.

Network and Host Portions of an IP Address

IP Address	172	.	16	.	4	.	1
	10101100		00010000		00000100		00000001
Subnet Mask	255	.	255	.	255	.	0
	11111111		11111111		11111111		00000000
Prefix /24 (24 high order bits)							

Alternative approach to Subnetting

Example

Divide network 192.168.1.0/24 into 4 subnets

Solution :

Old mask /24 → 11111111.11111111.11111111.00000000
255 . 255 . 255 . 0

4 subnets need 2 bits

- New subnet mask = 255 . 255 . 255 . 192
11111111.11111111.11111111.**11**000000 **(/26)**

- interesting octet is **192**

- hop count = $256 - 192 = 64$

- The first subnet is → 192.168.1.0/26
- The second subnet is → 192.168.1.64/26
- The third subnet is → 192.168.1.128/26
- The fourth subnet is → 192.168.1.192/26

Alternative approach to Subnetting

Example

Divide network 172.168.0.0/16 into 8 subnets

Solution :

Old mask /16 → 11111111.11111111.00000000.00000000

255 . 255 . 0 . 0

- 8 subnets need 3 bits

- new subnet mask = 255 . 255 . 224 . 0

11111111.11111111.**111**00000.00000000 (/19)

- interesting octet is **224**

- hop count = $256 - 224 = 32$

- The first subnet is → 172.168.0.0/19

- The second subnet is → 172.168.32.0/19

- The third subnet is → 172.168.64.0/19

-The 8th subnet is → 172.168.224.0/19

Alternative approach to Subnetting

Example

Divide network 10.0.0.0/10 into 4 subnets

Solution :

Old mask /10 → 11111111.11000000.00000000.00000000
255 . 192 . 0 . 0

4 subnets need 2 bits

- subnet mask = 255 . 240 . 0 . 0

11111111.11**11**1111.11111111.00000000 (/12)

- interesting octet is **240**

- hop count = 256 – 240 = **16**

- The first subnet is → 10.0.0.0/12

- The second subnet is → 10.16.0.0/12

- The third subnet is → 10.32.0.0/12

- The fourth subnet is → 10.48.0.0/12